



# 《计算机系统结构》课程直播

2020. 3.10

# 本次讲课： IEEE754 浮点数

1 形式

2 精度

3 运算与舍入

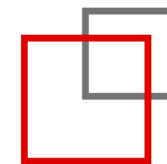
4 类型转换



# IEEE754 浮点数标准

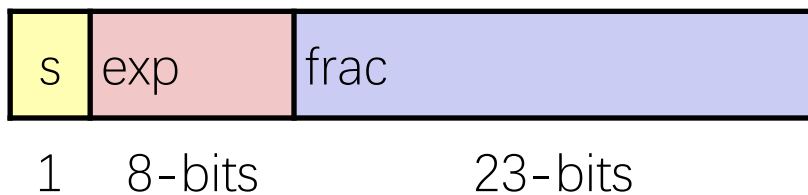


上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY



# 浮点数标准 IEEE754

- 单精度 Single precision: 32位



- 双精度 Double precision: 64 位



- 扩展精度 Extended precision: 80 位 (Intel)



# 单精度浮点数标准 IEEE754...

- 规格化数(Normal):



代表数值:  $(-1)^S \times 1.m \times 2^{e-\text{bias}}$

- Bias:
  - Single precision (8bits) : 127 (Exp: 1...254, E: -126...127)
  - Double precision (11-bits) : 1023 (Exp: 1...2046, E: -1022...1023)
- 规格化数的最高数字位总是1, IEEE754标准将这个1缺省存储(隐藏位), 使得尾数表示范围比实际存储多一位





# Example

Final representation:  $(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$

- Represent  $-0.75_{\text{ten}}$  in single and double-precision formats

Single:  $(1 + 8 + 23)$

Double:  $(1 + 11 + 52)$

-

# Example

Final representation:  $(-1)^S \times (1 + \text{Fraction}) \times 2^{(\text{Exponent} - \text{Bias})}$

- Represent  $-0.75_{\text{ten}}$  in single and double-precision formats

Single:  $(1 + 8 + 23)$

1 0111 1110 1000...000

Double:  $(1 + 11 + 52)$

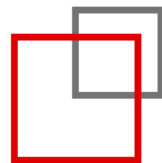
1 0111 1111 110 1000...000



# IEEE754浮点数精度



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY







# Floating Point (FP) Numbers

- **Floating point numbers:** numbers in scientific notation
  - Two uses
- Use I: real numbers (numbers with non-zero fractions)
  - 3.1415926...
  - 2.1878...
  - $6.62 * 10^{-34}$
- Use II: really big numbers
  - $6.02 * 10^{23}$
- Aside: best not used for currency values
  - Floating Point is Inexact, e.g. 0.1 (decimal)
  - `System.out.print("34.6-34.0=" + (34.6f-34.0f));`
  - `34.6-34.0=0.5999985`

# Floating Point is Inexact

- **Accuracy problems sometimes get bad**

- FP arithmetic not associative:  $(A+B)+C$  not same as  $A+(B+C)$
- Addition of big and small numbers
- summing many small numbers)
- Subtraction of two big numbers
- Example:  $(1*10^{30} + 1*10^0) - 1*10^{30} = (1*10^{30} - 1*10^{30}) = 0$
- In your code: **never test for equality between FP numbers**
  - Use something like: if  $(\text{abs}(a-b) < 0.00001)$  then ...

# IEEE 754 Standard Precision/Range

- **Single precision: float** in C
  - 32-bit: 1-bit sign + 8-bit exponent + 23-bit significand
  - Range:  $2.0 \times 10^{-38} < N < 2.0 \times 10^{38}$  ( 约为  $2^{127}$  )
  - Precision: 7 significant (decimal) digits (对应24位尾数位)
  - Used when exact precision is less important (e.g., 3D games)
- **Double precision: double** in C
  - 64-bit: 1-bit sign + 11-bit exponent + 52-bit significand
  - Range:  $2.0 \times 10^{-308} < N < 2.0 \times 10^{308}$  ( 约为  $2^{1023}$  )
  - Precision: 15 significant (decimal) digits (对应53位尾数位)
  - Used for scientific computations
- Numbers  $> 10^{308}$  don't come up in many calculations
  - $10^{80}$  ~ number of atoms in universe




# Exercise

---

- 设一个变量的值为4098，要求分别用32位补码整数和IEEE 754单精度浮点格式表示该变量（结果用十六进制表示），并说明哪段二进制序列在两种表示中完全相同，为什么会相同？

# Exercise

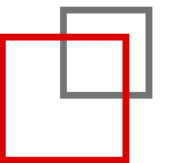
- 设一个变量的值为 $-2147483647$ ，要求分别用32位补码整数和IEEE754单精度浮点格式表示该变量（结果用十六进制表示），并说明哪种表示其值完全精确，哪种表示的是近似值。



# IEEE754浮点数运算与舍入



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY





# FP Addition – Binary Example

- Consider the following binary example

$$1.010 \times 2^1 + 1.100 \times 2^3$$

Convert to the larger exponent:

$$0.0101 \times 2^3 + 1.1000 \times 2^3$$

Add

$$1.1101 \times 2^3$$

Normalize

$$1.1101 \times 2^3$$

Check for overflow/underflow

Round

Re-normalize

IEEE 754 format: 0 10000010 110100000000000000000000



# FP Multiplication

---

- Similar steps:
  - Compute exponent (careful!)
  - Multiply significands (set the binary point correctly)
  - Normalize
  - Round (potentially re-normalize)
  - Assign sign

# 讨论



对于IEEE754单精度浮点数加减运算，只要对阶时得到的两个阶码之差的绝对值 $|E_1 - E_2|$ 大于等于（），就无须继续进行后续处理，此时运算结果直接取阶大的那个数

- ☐ A. 24
- ☒ B. 25
- ☐ C. 126
- ☐ D. 128

正确答案：B 你选对了

## 二进制数的最近舍入



- 二进制数的 Round-To-Even
  - 偶数 “Even”：最低有效位为 0
  - 数值处于中间 “Half way” 时，要舍弃的位数的形式  $= 100\cdots_2$

- 举例：Round to nearest  $1/4$  (2 bits right of binary point)

Value	Binary	Rounded	Action	Rounded Value
$2 \frac{3}{32}$	$10.00\underline{011}_2$	$10.00_2$	( $<1/2$ —down)	2
$2 \frac{3}{16}$	$10.00\underline{110}_2$	$10.01_2$	( $>1/2$ —up)	$2 \frac{1}{4}$
$2 \frac{7}{8}$	$10.11\underline{100}_2$	$11.00_2$	( $=1/2$ —up)	3
$2 \frac{5}{8}$	$10.10\underline{100}_2$	$10.10_2$	( $=1/2$ —down)	$2 \frac{1}{2}$

# Exercise



- Given  $A=2.6125 \times 10^1$ ,  $B=4.150390625 \times 10^{-1}$ ,
- Calculate the sum of A and B by hand, assuming A and B are stored by the following format,
- Assume 1 guard(保护位), 1 round bit (舍入位), and 1 sticky bit (粘滞位) and round to the nearest even (首选“偶数”值舍入) .
- IEEE754规定, 浮点运算的中间结果的右边都必须额外多保留两位 (保护位、舍入位) 为获得无限精度求出后的舍入效果, 再加一个粘滞位。

- $2.6125 \times 10^1 + 4.150390625 \times 10^{-1}$

$$2.6125 \times 10^1 = 26.125 = 11010.001 = 1.1010001000 \times 2^4$$

$$4.150390625 \times 10^{-1} = .4150390625 = .011010100111$$

$$= 1.1010100111 \times 2^{-2} \quad (\text{对阶, 小阶往大阶对})$$

Shift binary point 6 to the left to align exponents,

GR

$$1.1010001000 \ 00$$

$$+ .\textcolor{red}{000001}1010 \ 10 \ 0111 \ (\text{Guard} = 1, \text{Round} = 0, \text{Sticky} = 1)$$

-----

$$1.1010100010 \ 10 \quad (\text{尾数相加}) \quad \text{and} \quad (\text{尾数规格化检查})$$

the extra bits (G,R,S) are more than half of the least significant bit (0).

Thus, the value is rounded up. (舍入)

$$1.1010100011 \times 2^4 \quad (\text{检查, 无溢出})$$

$$= 11010.100011 \times 2^0 = 26.546875 = 2.6546875 \times 10^1$$



# Rounding

1 . BBBG**RXXX**

Guard bit: 1<sup>st</sup> bit removed

Round bit: 2<sup>nd</sup> bit removed

Sticky bit:  
OR of remaining bits

- Round up conditions

<i>Value</i>	<i>Fraction</i>	<i>GRS</i>	<i>Incr?</i>	<i>Rounded</i>
128	1.000 <b>0000</b>	000	N	1.000
15	1.101 <b>0000</b>	000	N	1.101
17	1.000 <b>1000</b>	100	N	1.000
19	1.001 <b>1000</b>	100	Y	1.010
138	1.000 <b>1010</b>	101	Y	1.001
63	1.111 <b>1100</b>	110	Y	10.000

# Postnormalize

- Issue
  - Rounding may have caused overflow
  - Handle by shifting right once & incrementing exponent

<i><b>Value</b></i>	<i><b>Rounded</b></i>	<i><b>Exp</b></i>	<i><b>Adjusted</b></i>	<i><b>Result</b></i>
128	1.000	7		128
15	1.101	3		15
17	1.000	4		16
19	1.010	4		20
138	1.001	7		134
63	10.000	5	1.000/6	64

# Arithmetic Latencies

- Latency in cycles of common arithmetic operations
- Source: *Agner Fog*, <https://www.agner.org/optimize/#manuals>
  - AMD Ryzen core

	Int 32	Int 64	Fp 32	Fp 64
Add/Subtract	1	1	5	5
Multiply	3	3	5	5
Divide	14-30	14-46	8-15	8-15

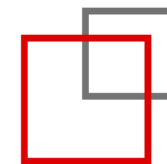
- Divide is variable latency based on the size of the dividend
  - Detect number of leading zeros, then divide
- Why is FP divide faster than integer divide?



# 浮点数类型转换



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY



# 练习



设int x=1,float y=2,则表达式x/y的值是: ()

- A. 0
- B. 1
- C. 2
- D. 以上都不是

# Floating Point in C



- C Guarantees Two Levels
  - **float**      single precision
  - **double**     double precision
- Conversions/Casting
  - Casting between **int**, **float**, and **double** changes bit representation
  - **double/float** → **int**
    - Truncates fractional part
    - Like rounding toward zero
    - Not defined when out of range or NaN: Generally sets to TMin
  - **int** → **double**
    - Exact conversion, as long as **int** has  $\leq 53$  bit word size
  - **int** → **float**
    - Will round according to rounding mode



# Floating Point Puzzles



- For each of the following C expressions, either:
  - Argue that it is true for all argument values
  - Explain why not true

```
int x = ...;  
float f = ...;  
double d = ...;
```

Assume neither  
**d** nor **f** is NaN

- `x == (int)(float) x`
- `x == (int)(double) x`
- `f == (float)(double) f`
- `d == (double)(float) d`
- `f == -(-f);`
- `2/3 == 2/3.0`
- `d < 0.0`  $\Rightarrow$  `((d*2) < 0.0)`
- `d > f`  $\Rightarrow$  `-f > -d`
- `d * d >= 0.0`
- `(d+f) - d == f`

# Interesting Numbers

**{single, double}**



<i>Description</i>	<i>exp</i>	<i>frac</i>	<i>Numeric Value</i>
▪ Zero	00...00	00...00	0.0
▪ Smallest Pos. Denorm. <ul style="list-style-type: none"><li>▪ Single <math>\approx 1.4 \times 10^{-45}</math></li><li>▪ Double <math>\approx 4.9 \times 10^{-324}</math></li></ul>	00...00	00...01	$2^{-\{23,52\}} \times 2^{-\{126,1022\}}$
▪ Largest Denormalized <ul style="list-style-type: none"><li>▪ Single <math>\approx 1.18 \times 10^{-38}</math></li><li>▪ Double <math>\approx 2.2 \times 10^{-308}</math></li></ul>	00...00	11...11	$(1.0 - \epsilon) \times 2^{-\{126,1022\}}$
▪ Smallest Pos. Normalized <ul style="list-style-type: none"><li>▪ Just larger than largest denormalized</li></ul>	00...01	00...00	$1.0 \times 2^{-\{126,1022\}}$
▪ One	01...11	00...00	1.0
▪ Largest Normalized <ul style="list-style-type: none"><li>▪ Single <math>\approx 3.4 \times 10^{38}</math></li><li>▪ Double <math>\approx 1.8 \times 10^{308}</math></li></ul>	11...10	11...11	$(2.0 - \epsilon) \times 2^{\{127,1023\}}$

# Special Properties of the IEEE Encoding



- FP Zero Same as Integer Zero
  - All bits = 0
- Can (Almost) Use Unsigned Integer Comparison
  - Must first compare sign bits
  - Must consider  $-0 = 0$
  - NaNs problematic
    - Will be greater than any other values
    - What should comparison yield?
  - Otherwise OK
    - Denorm vs. normalized
    - Normalized vs. infinity

# Summary



- IEEE Floating Point has clear mathematical properties
- Represents numbers of form  $M \times 2^E$
- One can reason about operations independent of implementation
  - As if computed with perfect precision and then rounded
- Not the same as real arithmetic
  - Violates associativity/distributivity
  - Makes life difficult for compilers & serious numerical applications programmers



## 下一节

- 周四 8: 00
- 存储系统
- 请做好准备

# 再见

