



# 《高级计算机系统结构》课程直播

## 2020. 5.9

听不到声音请及时调试声音设备，可以下课后补签到

请将ZOOM名称改为“姓名”；

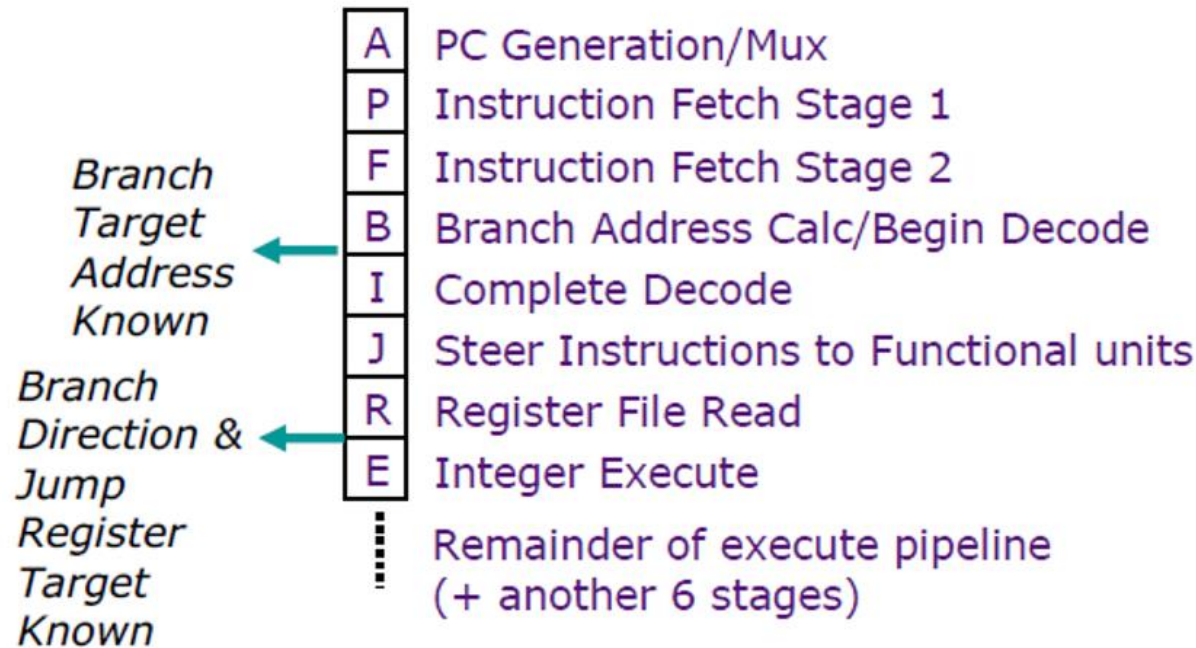
# 本节内容

---

- 通知：第三次作业发布，截止期**5.17日**
- 习题讲解
  - 转移预测
  - 多线程
- 数据级并行性
  - Vector, SIMD Architectures、GPU

From : H&P Computer Architecture: A Quantitative Approach,  
Fifth Edition, (5th edition)

# 习题1： Pipeline of UltraSparc-III processor



Instruction	Taken known? (At the end of)	Target known? (At the end of)
BEQZ/BNEZ	R	B
J	B (always taken)	B
JR	B (always taken)	R

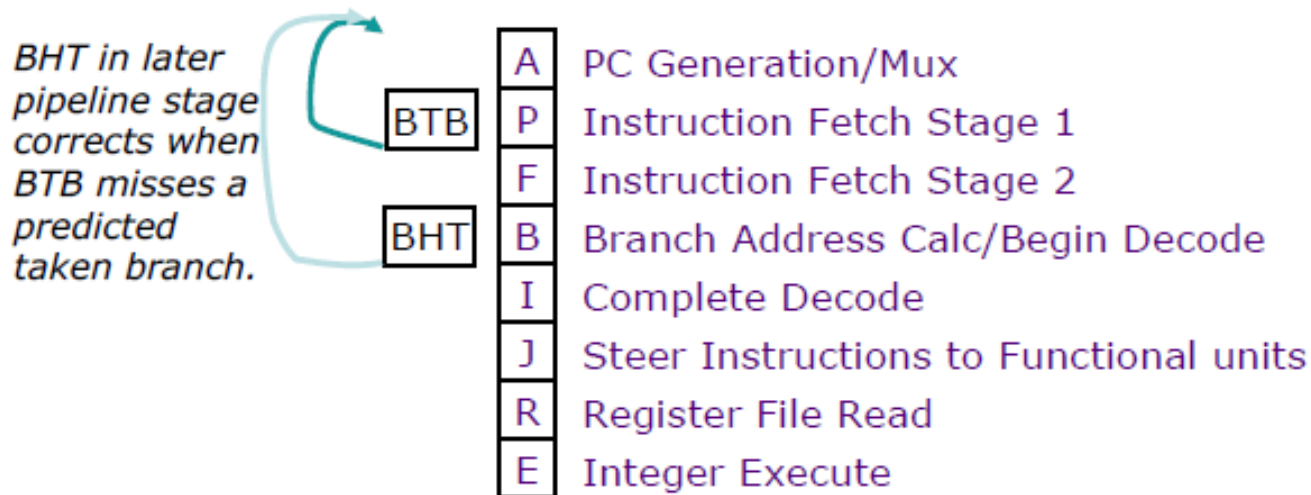
# BHT 的作用

- ❑ First, we add a branch history table (BHT) in the fetch pipeline
- ❑ In the B stage (Branch Address Calc/Begin Decode), a conditional branch instruction (BEQZ/BNEZ) looks up the BHT, but an unconditional jump does not. If a branch is predicted to be taken
- ❑ some of the instructions are flushed and the PC is redirected to the calculated branch target address.
- ❑ The instruction at PC+4 is fetched by default unless PC is redirected by an older instruction.

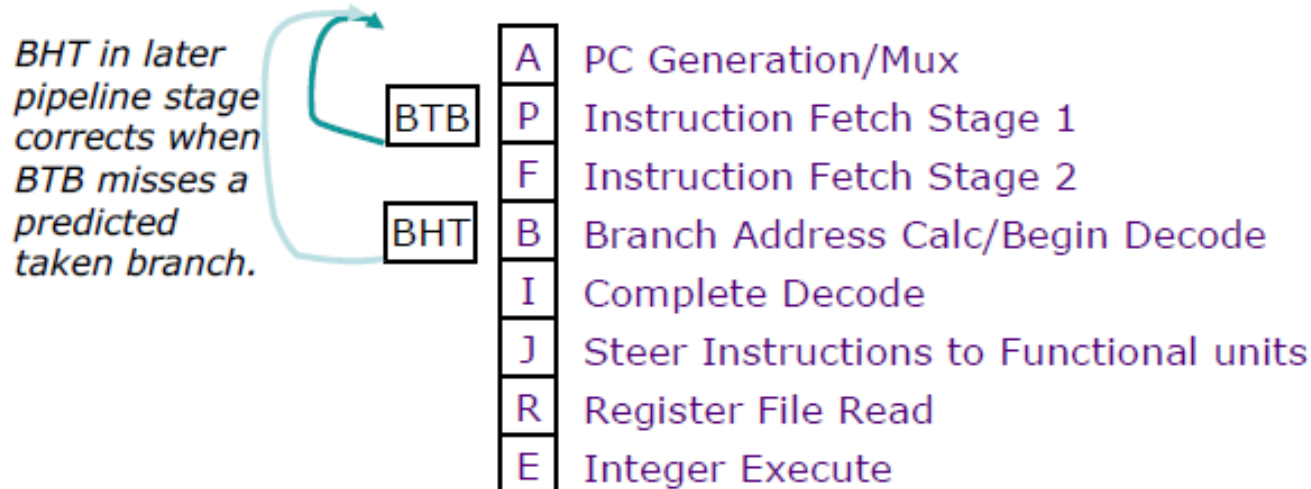
BHT	A	PC Generation/Mux		Predicted Taken?	Actually Taken?	Pipeline bubbles
	P	Instruction Fetch Stage 1	BEQZ/ BNEZ	Y	Y	3
	F	Instruction Fetch Stage 2		Y	N	6
	B	Branch Address Calc/Begin Decode		N	Y	6
	I	Complete Decode		N	N	0
	J	Steer Instructions to Functional	J	Always taken (No lookup)	Y	3
	R	Register File Read	JR	Always taken (No lookup)	Y	6
	E	Integer Execute				

# BTB的作用

- ❑ 在取指阶段增加一个转移目标缓冲器BTB
- ❑ BTB表中,会为预测为“转移”的jump或branch 记录一对表项(当前指令的PC, 转移目标指令的PC). 假设在BTB表中, 转移目标即target\_PC一直是正确的(当然, 转移方向的预测仍然有可能是错误的);
- ❑ The BTB is looked up every cycle. If there is a match with the current PC, PC is redirected to the target\_PC predicted by the BTB (unless PC is redirected by an older instruction); if not, it is set to PC+4.



# BTB的作用



	BTB Hit?	(BHT) Predicted Taken?	Actually Taken?	Pipeline bubbles
Conditional Branches	Y	Y	Y	1
	Y	Y	N	6
	Y	N	Y	Cannot occur
	Y	N	N	Cannot occur
	N	Y	Y	3
	N	Y	N	6
	N	N	Y	6
	N	N	N	0



## 习题2：多线程

---

- ❑ This problem evaluates the effectiveness of multithreading using a simple database benchmark.
- ❑ The benchmark searches for an entry in a linked list built from the following structure, which contains a key, a pointer to the next node in the linked list, and a pointer to the data entry.

```
struct node {  
    int key;  
    struct node *next;  
    struct data *ptr;  
}
```

## 习题2：多线程（续）

- ❑ The following MIPS code shows the core of the benchmark, which traverses the linked list and finds an entry with a particular key. Assume MIPS has no delay slots.

; R1: a pointer to the linked list

; R2: the key to find

loop: LW R3, 0(R1) ; load a key

LW R4, 4(R1) ; load the next pointer

SEQ R3, R3, R2 ; set R3 if R3 == R2

BNEZ R3, End ; found the entry

ADD R1, R0, R4

BNEZ R1, Loop ; check the next node

End: ; R1 contains a pointer to the matching entry or zero if not found

- ❑ We run this benchmark on a single-issue in-order processor.
- ❑ If an instruction cannot be issued due to a data dependency, the processor stalls.
- ❑ Integer instructions take one cycle to execute and the result can be used in the next cycle.
- ❑ if SEQ is executed in cycle 1, BNEZ can be executed in cycle 2.
- ❑ We also assume that the processor has a perfect branch predictor with no penalty for both taken and not-taken branches.



## the loop takes 104 cycles to execute

- ❑ Assume that our system does not have a cache. Each memory operation directly accesses main memory and takes 100 CPU cycles. The load/store unit is fully pipelined, and nonblocking. After the processor issues a memory operation, it can continue executing instructions until it reaches an instruction that is dependent on an outstanding memory operation.
- ❑ The following table summarizes the execution time for each instruction. From the table, the loop takes 104 cycles to execute.

Instruction		Start Cycle	End Cycle
LW	R3, 0(R1)	1	100
LW	R4, 4(R1)	2	101
SEQ	R3, R3, R2	101	101
BNEZ	R3, End	102	102
ADD	R1, R0, R4	103	103
BNEZ	R1, Loop	104	104

## Problem A.


---

- ❑ Now we add zero-overhead multithreading to our pipeline.
- ❑ A processor executes multiple threads, each of which performs an independent search. Hardware mechanisms schedule a thread to execute each cycle.
- ❑ In our first implementation, the processor switches to a different thread every cycle using fixed round robin scheduling. Each of the  $N$  threads executes one instruction every  $N$  cycles.
- ❑ What is the minimum number of threads that we need to fully utilize the processor, i.e., execute one instruction per cycle?


- 
- ❑ How does multithreading affect throughput (number of keys the processor can find within a given time) and latency (time processor takes to find an entry with a specific key)?
  - ❑ Assume the processor switches to a different thread every cycle and is fully utilized.

- 
- ❑ We change the processor to only switch to a different thread when an instruction cannot execute due to data dependency.
  - ❑ What is the minimum number of threads to fully utilize the processor now?
  - ❑ Note that the processor issues instructions in-order in each thread.

## 粗粒度多线程: Coarse-Grain Multithreading




CGMT




## 细粒度多线程: Fine-Grain Multithreading

- FGMT
  - **Multiple threads in pipeline at once**
  - (Many) more threads




## 同时多线程: Simultaneous Multithreading



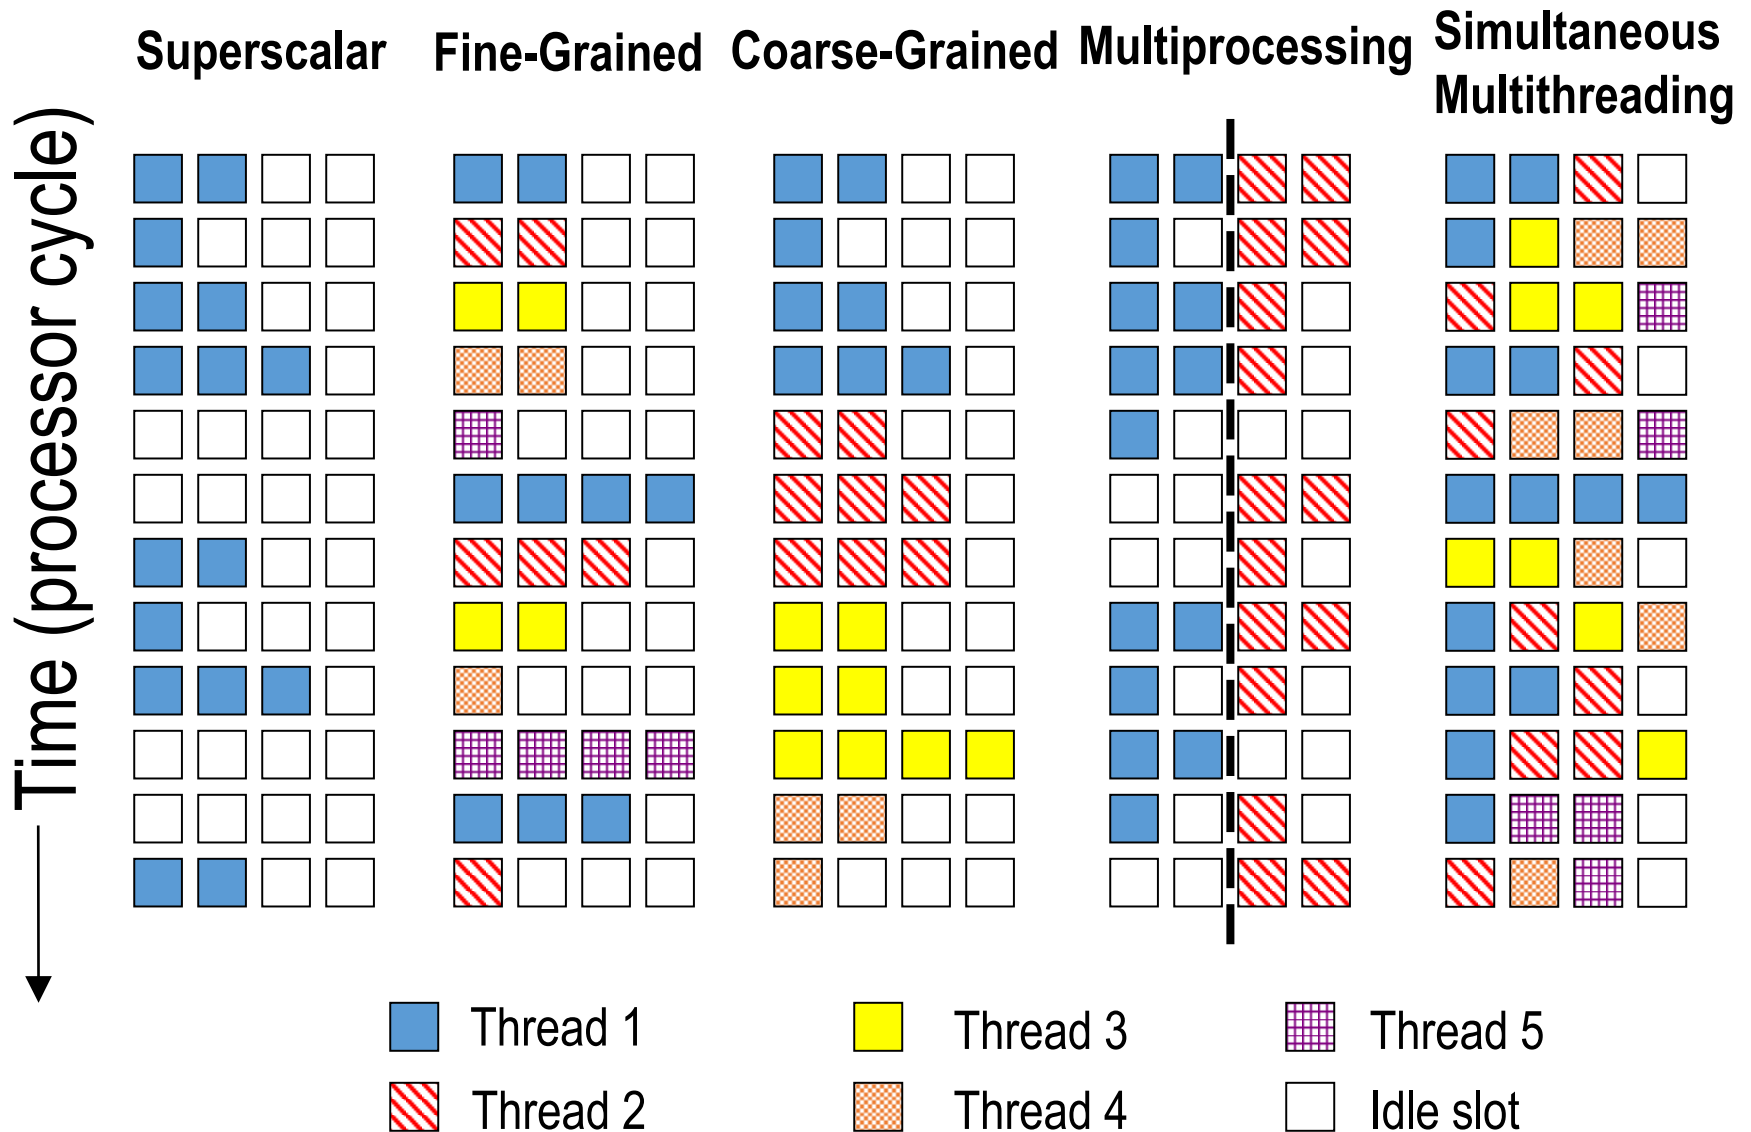
- SMT

- Replicate map table, share (larger) physical register file





# Summary: Multithreaded Categories



---

# 数据级并行性

Data Level Parallelism

# Data-Level Parallelism

---

- ❑ 数据级并行的研究动机
  - 传统指令级并行技术的问题
  - SIMD结构的优势
  - 数据级并行的种类
- ❑ 向量体系结构
  - 向量处理模型
  - 起源-超级计算机
  - 基本特性及结构
  - 性能评估及优化
- ❑ 面向多媒体应用的SIMD指令集扩展
- ❑ GPU
  - GPU简介
  - GPU的编程模型
  - GPU的存储系统

# SIMD 结构的种类

---

## ❑ 向量体系结构

## ❑ 多媒体SIMD指令集 扩展

- x86 processors:

- ❑ 每年增加2cores/chip

- ❑ SIMD 宽度每4年翻一番

- ❑ SIMD潜在加速比是MIMD的2倍

## ❑ Graphics Processor Units (GPUs)

# Vector Code Element-by-Element Multiplication

---

## # C code

```
for (i=0; i<64; i++)  
    C[i] = A[i] * B[i];
```

## # Scalar Assembly Code

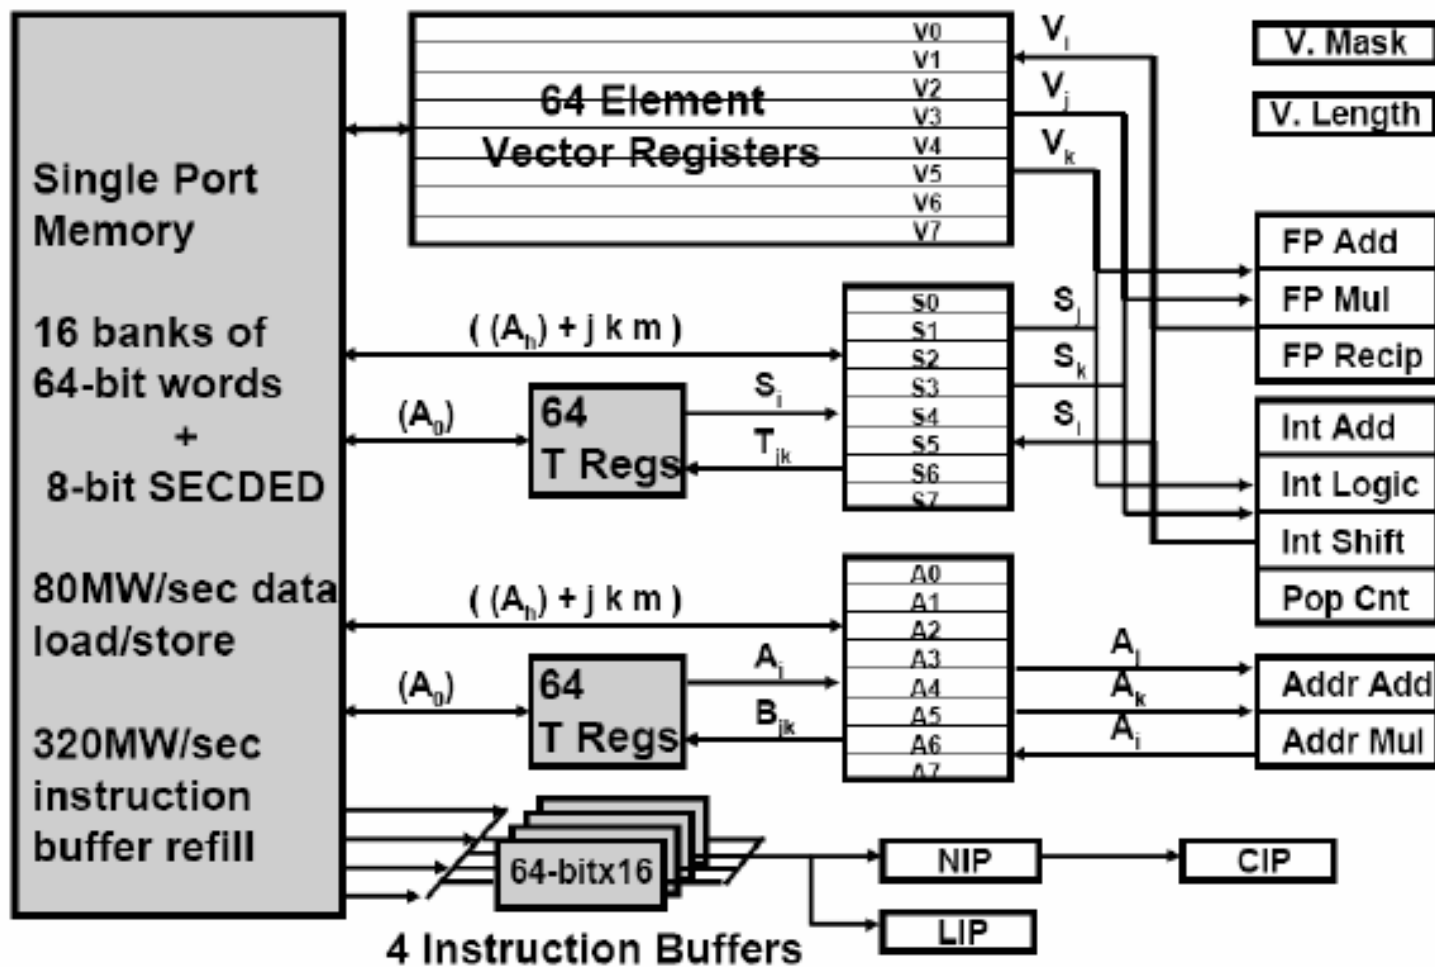
```
    LI R4, 64  
loop:  
    L.D F0, 0(R1)  
    L.D F2, 0(R2)  
    MUL.D F4, F2, F0  
    S.D F4, 0(R3)  
    DADDIU R1, 8  
    DADDIU R2, 8  
    DADDIU R3, 8  
    DSUBIU R4, 1  
    BNEZ R4, loop
```

## # Vector Assembly Code

```
    LI VLR, 64  
    LV V1, R1  
    LV V2, R2  
    MULVV.D V3, V1, V2  
    SV V3, R3
```

# 向量处理机的基本组成单元

## Cray-1 (1976)



内存体读写：50ns； 处理器周期：12.5ns（80MHz）

# 向量体系结构

---

## □ 向量处理机**基本概念**

- 基本思想：两个向量的对应分量进行运算，产生一个结果向量

## □ 向量处理机**基本特征**

- VSIW- 一条指令包含多个操作
- 单条向量指令内所包含的操作相互独立
- 以已知模式访问存储器-多体交叉存储系统
- 控制相关少

## □ 向量处理机**基本结构**

- 向量指令并行执行
- 向量运算部件的执行方式-流水线方式
- 向量部件结构-多“道”结构-多条运算流水线


## □ 向量处理机**性能优化**

- 链接技术
- 条件执行



# Vector Arithmetic Execution

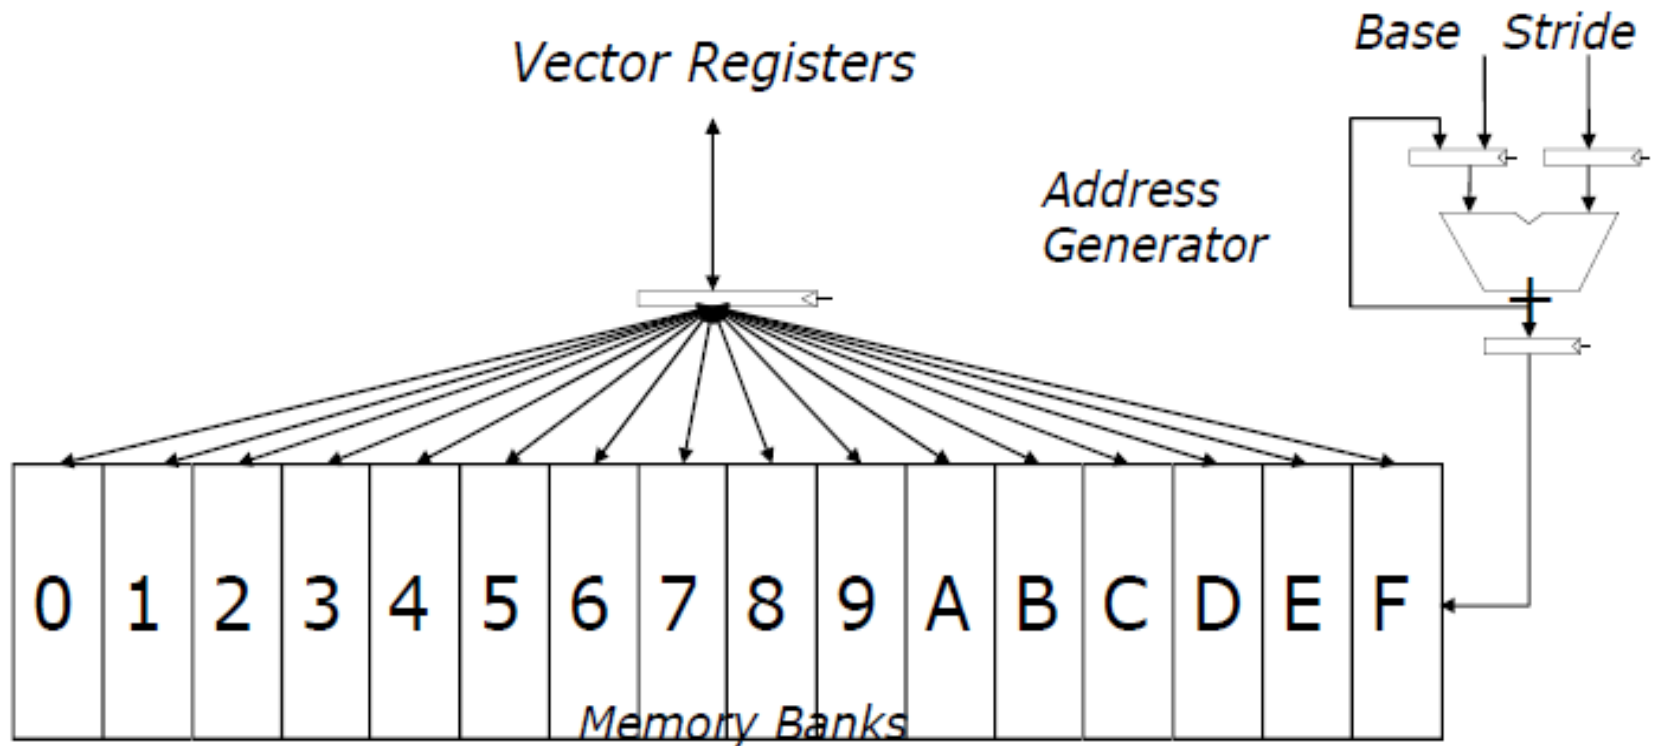
- ❑ Use deep pipeline ( $\Rightarrow$  fast clock) to execute element operations
- ❑ Simplifies control of deep pipeline because elements in vector are independent
  - no data hazards!
  - no bypassing needed



# Interleaved Vector Memory System


Cray-1, 16 banks, 4 cycle bank busy time, 12 cycle latency

- *Bank busy time*: Time before bank ready to accept next request



# Vector Chaining

- ❑ Vector version of register bypassing
  - introduced with Cray-1



# Vector Conditional Execution

- ❑ Problem: Want to vectorize loops with conditional code:

```
for (i=0; i<N; i++)  
    if (A[i]>0) then A[i] = B[i];
```

- ❑ Solution: Add vector mask (or flag) registers
- ❑ Code example:

```
CVM                # Turn on all elements  
LV VA, RA          # Load entire A vector  
SGTVS.D VA, F0     # Set bits in mask register where A>0  
LV VA, RB          # Load B vector into A under mask  
SV VA, RA          # Store A back to memory under mask
```

# Masked Vector Instructions

## Simple Implementation

- execute all N operations, turn off result writeback according to mask


M[7]=1 A[7] B[7]

M[6]=0 A[6] B[6]

M[5]=1 A[5] B[5]

M[4]=1 A[4] B[4]

M[3]=0 A[3] B[3]



## Density-Time Implementation

- scan mask vector and only execute elements with non-zero masks

M[7]=1

M[6]=0

M[5]=1

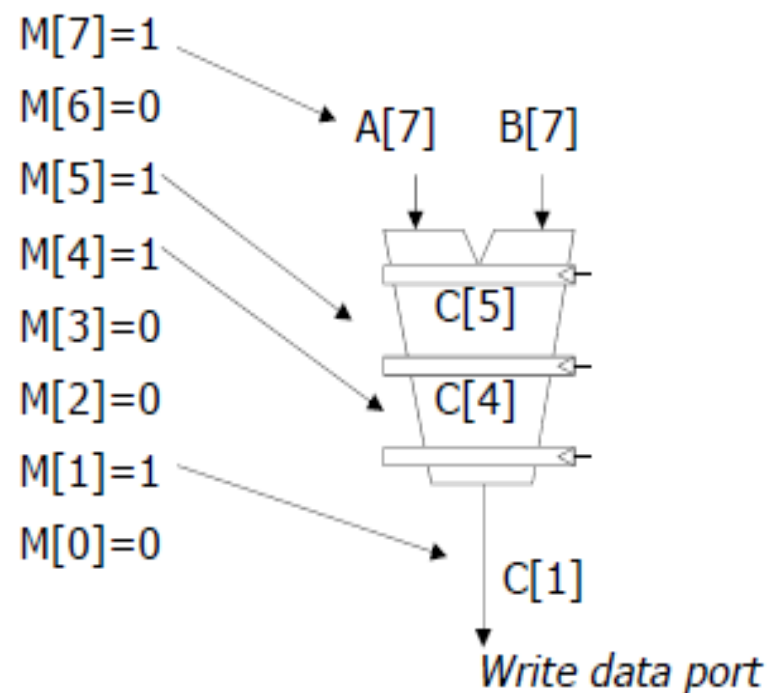
M[4]=1

M[3]=0

M[2]=0

M[1]=1

M[0]=0



# Example Vector Machines

Machine	Year	Clock(MHz)	Regs	Elements	Fus	LSUs
Cray 1	1976	80	8	64	6	1
Cray XMP	1983	120	8	64	8	2L, 1S
Cray YMP	1988	166	8	64	8	2L, 1S
Cray C-90	1991	240	8	128	8	4
Cray T-90	1996	455	8	128	8	4
Conv. C-1	1984	10	8	128	4	1
Conv. C-4	1994	133	16	128	3	1
Fuj. VP200	1982	133	8-256	32-1024	3	2
Fuj. VP300	1996	100	8-256	32-1024	3	2
NEC SX/2	1984	160	8+8K	256+var	16	8
NEC SX/3	1995	400	8+8K	256+var	16	8

# Array vs. Vector Processors

**Array processor:** 又称为并行处理机、SIMD处理器。其核心是一个由**多个处理单元**构成的**阵列**，用**单一的控制部件**来控制多个处理单元对**各自的数据**进行**相同的运算**和操作。

ARRAY PROCESSOR

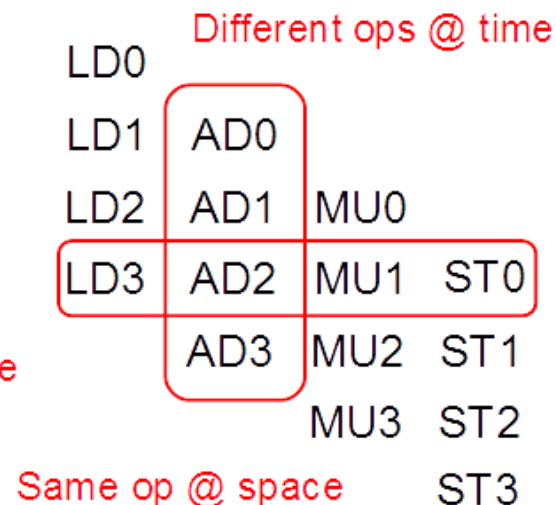
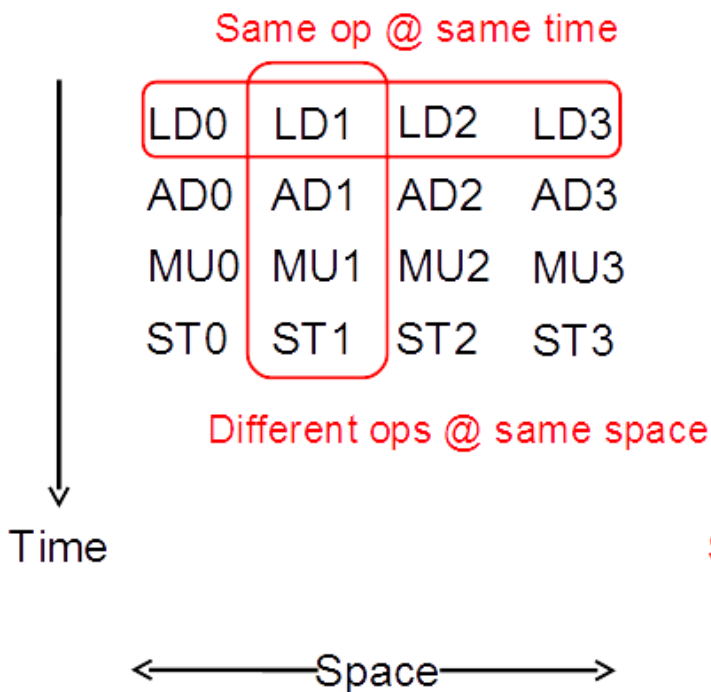


VECTOR PROCESSOR



Instruction Stream

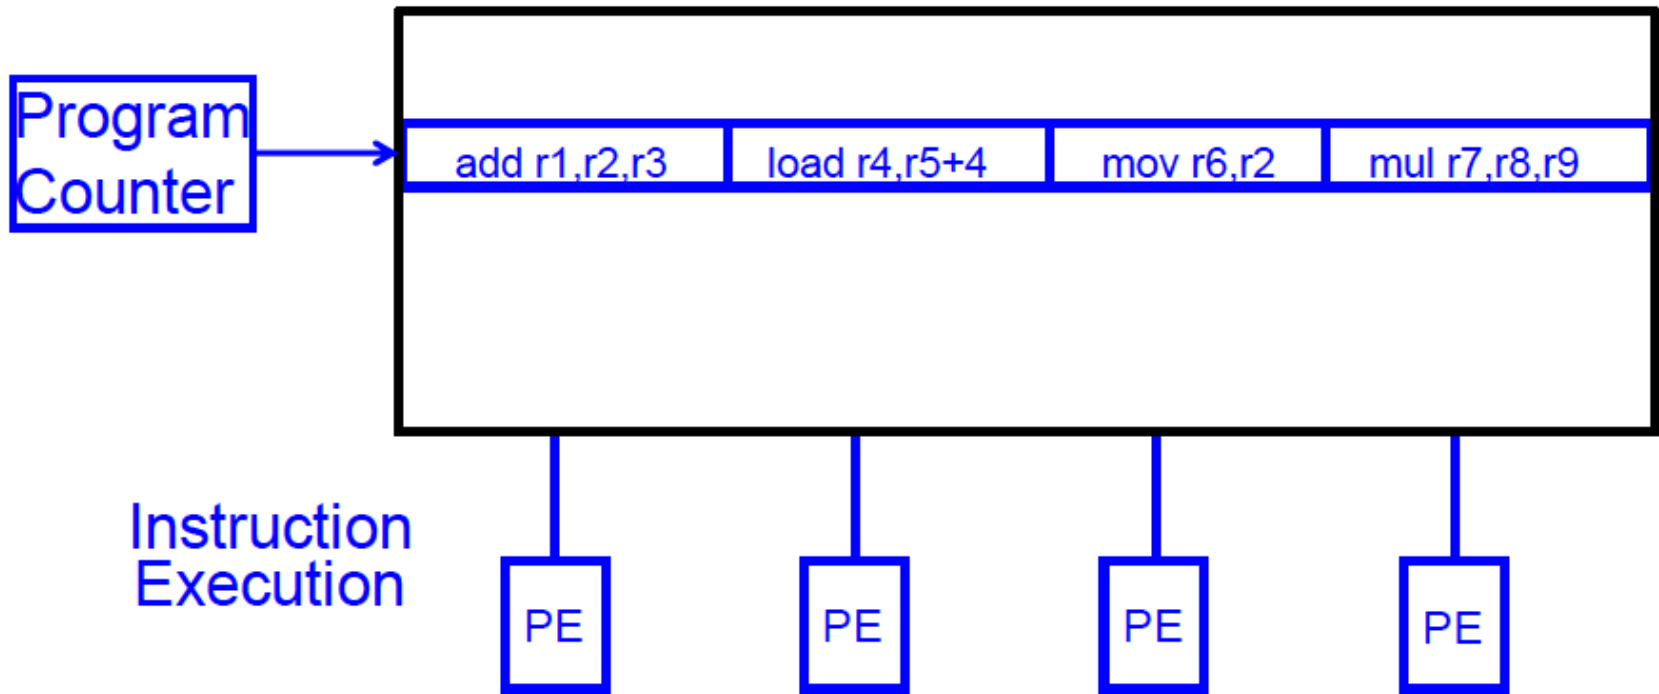
```
LD   VR ← A[3:0]
ADD  VR ← VR, 1
MUL  VR ← VR, 2
ST   A[3:0] ← VR
```





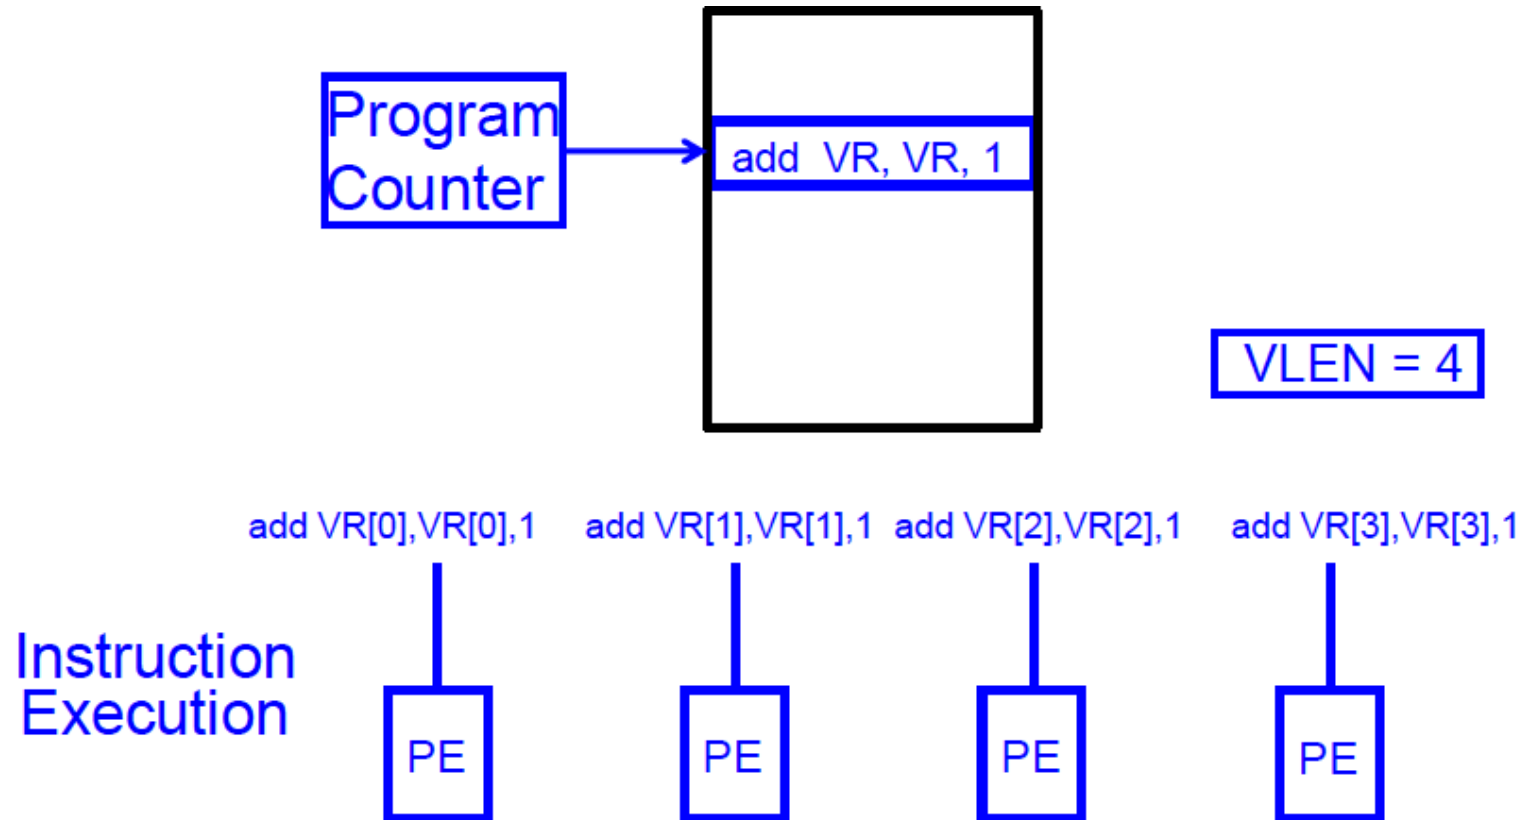
# SIMD Array Processing vs. VLIW

- ❑ VLIW: 多个独立的操作由编译器封装在一起



# SIMD Array Processing vs. VLIW

- ❑ Array processor: 单个操作作用在多个不同的数据元素上



---


# **Data-Level Parallelism in SIMD extensions**

# Multimedia Extensions (aka SIMD extensions)

- ❑ 在已有的ISA中添加一些向量长度很短的向量操作指令
- ❑ 将已有的 64-bit 寄存器拆分为 2x32b or 4x16b or 8x8b
  - 1957年, Lincoln Labs TX-2 将36bit datapath 拆分为 2x18b or 4x9b
  - 新的设计具有较宽的寄存器
    - 📁 128b for PowerPC AltiVec, Intel SSE2/3/4
    - 📁 256b for Intel AVX (Advanced Vector Extensions)
- ❑ 单条指令可实现寄存器中所有向量元素的操作

# Intel Pentium MMX Operations


- ❑ idea: 一条指令操作同时作用于不同的数据元
  - 全阵列处理
  - 用于多媒体操作



(a)



(b)



(c)



(d)

- No VLEN register
- Opcode determines data type:
  - 8 8-bit bytes
  - 4 16-bit words
  - 2 32-bit doublewords
  - 1 64-bit quadword
- Stride always equal to 1.

Figure 1. MMX technology data types: packed byte (a), packed word (b), packed doubleword (c), and quadword (d).

# Example SIMD Code

## • Example DXPY:

L.D	F0,a	;load scalar a
MOV	F1, F0	;copy a into F1 for SIMD MUL
MOV	F2, F0	;copy a into F2 for SIMD MUL
MOV	F3, F0	;copy a into F3 for SIMD MUL
DADDIU	R4,Rx,#512	;last address to load
Loop:	L.4D F4,0[Rx]	;load X[i], X[i+1], X[i+2], X[i+3]
MUL.4D	F4,F4,F0	$;a \times X[i], a \times X[i+1], a \times X[i+2], a \times X[i+3]$
L.4D	F8,0[Ry]	;load Y[i], Y[i+1], Y[i+2], Y[i+3]
ADD.4D	F8,F8,F4	$;a \times X[i] + Y[i], \dots, a \times X[i+3] + Y[i+3]$
S.4D	0[Ry],F8	;store into Y[i], Y[i+1], Y[i+2], Y[i+3]
DADDIU	Rx,Rx,#32	;increment index to X
DADDIU	Ry,Ry,#32	;increment index to Y
DSUBU	R20,R4,Rx	;compute bound
BNEZ	R20,Loop	;check if done

# Multimedia Extensions versus Vectors

## ❑ 受限的指令集:

- 无向量长度控制
- Load/store操作无 常数步长寻址和 scatter/gather操作
- loads 操作必须64/128-bit 边界对齐

## ❑ 受限的向量寄存器长度:

- 需要超标量发射以保持multiply/add/load 部件忙
- 通过循环展开隐藏延迟增加了寄存器读写压力

## ❑ 在微处理器设计中向全向量化发展

- 更好地支持非对齐存储器访问
- 支持双精度浮点数操作 (64-bit floating-point)
- Intel AVX spec (announced April 2008), 256b vector registers (expandable up to 1024b)



## 下一节

---

### □ 数据级并行性

- GPU

From : H&P Computer Architecture: A Quantitative Approach,  
Fifth Edition, (5th edition)



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY



谢 谢！

