



# 《计算机系统结构》课程直播

## 2020. 5.19

听不到声音请及时调试声音设备，可以下课后补签到

请将ZOOM名称改为“姓名”；

# 本节内容

---

## □ 数据级并行性

- GPU存储模型
- GPU的性能

## □ 线程级并行性

- 高速缓存一致性

From : H&P Computer Architecture: A Quantitative Approach,  
Fifth Edition, (5th edition)

---

# Graphics Processing Units

GPU

# Using CPU+GPU Architecture

## ❑ CPU+GPU异构多核系统

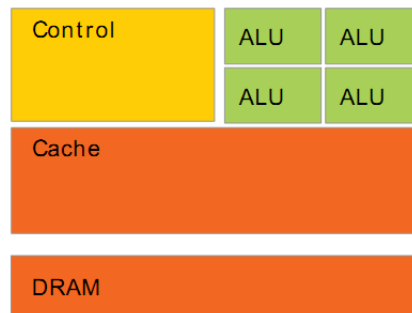
- 针对每个任务选择合适的处理器和存储器

## ❑ 通用CPU 适合执行一些串行的线程

- 串行执行快
- 带有cache，访问存储器延时低


## ❑ GPU 适合执行大量并行线程

- 可扩放的并行执行
- 高带宽的并行存取



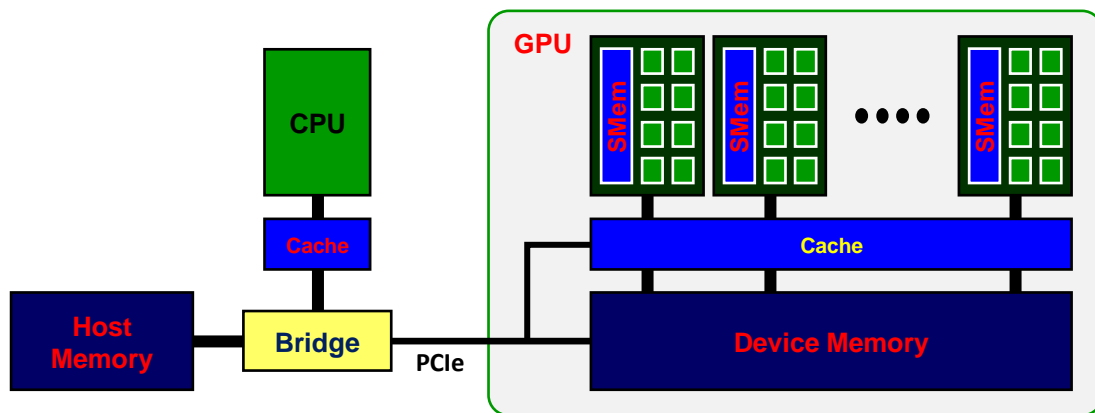
CPU

强控制、弱  
计算



GPU

弱控制、强  
计算



# Heterogeneous Computing

```
#include <iostream>
#include <algorithm>

using namespace std;

#define N 1024
#define RADIUS 3
#define BLOCK_SIZE 16

__global__ void stencil_1d(int *in, int *out) {
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;
    int lindex = threadIdx.x + RADIUS;

    // Read input elements into shared memory
    temp[lindex] = in[gindex];
    if (threadIdx.x < RADIUS) {
        temp[lindex - RADIUS] = in[gindex - RADIUS];
        temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];
    }

    // Synchronize (ensure all the data is available)
    __syncthreads();

    // Apply the stencil
    int result = 0;
    for (int offset = -RADIUS; offset <= RADIUS; offset++)
        result += temp[lindex + offset];

    // Store the result
    out[gindex] = result;
}

void fill_ints(int *x, int n) {
    fill_n(x, n, 1);
}

int main(void) {
    int *in, *out; // host copies of a, b, c
    int *d_in, *d_out; // device copies of a, b, c
    int size = (N + 2 * RADIUS) * sizeof(int);

    // Alloc space for host copies and setup values
    in = (int *)malloc(size); fill_ints(in, N + 2 * RADIUS);
    out = (int *)malloc(size); fill_ints(out, N + 2 * RADIUS);

    // Alloc space for device copies
    cudaMalloc((void **)&d_in, size);
    cudaMalloc((void **)&d_out, size);

    // Copy to device
    cudaMemcpy(d_in, in, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_out, out, size, cudaMemcpyHostToDevice);

    // Launch stencil_1d() kernel on GPU
    stencil_1d<<<N/BLOCK_SIZE, BLOCK_SIZE>>>>(d_in + RADIUS,
    d_out + RADIUS);

    // Copy result back to host
    cudaMemcpy(out, d_out, size, cudaMemcpyDeviceToHost);


    // Cleanup
    free(in); free(out);
    cudaFree(d_in); cudaFree(d_out);
    return 0;
}
```

parallel fn

serial code

parallel code

serial code



# GPU: a multithreaded coprocessor


SP: scalar processor  
'CUDA core'

Executes one thread

**SM**

streaming multiprocessor  
32xSP (or 16, 48 or more)

Fast local '**shared memory**'  
(shared between SPs)  
16 KiB (or 64 KiB)



# 总结：GPU的三个key idea

## □ Employ multiple processing cores

- 简单的核 (比起ILP， 更注重TLP：thread-level parallelism)

## □ Pack cores full of ALU

- 开发数据级的并行性 (SIMD)
- 一条指令执行时，多个ALU同步处理不同的数据
- 芯片面积额外增加的开销不大，但大大增强了计算能力

## □ 利用多线程提高系统的处理能力（吞吐量）

- 轮流交替执行不同线程的代码段以隐藏延迟

# Nvidia 通用图形加速单元体系结构

- ❑ 2008 Tesla
  - ❑ 2010 Fermi
  - ❑ 2012 Kepler
  - ❑ 2014 Maxwell
  - ❑ 2016 Pascal
- 
- ❑ 每个SM包含的SP（GPU core）数量依据GPU架构而不同，Fermi架构GF100是32个，GF10X是48个，Kepler架构都是192个，Maxwell都是128个，Pascal64个。



# Floorplan of Fermi GTX480



# Pascal 架构(Tesla P100 2016)

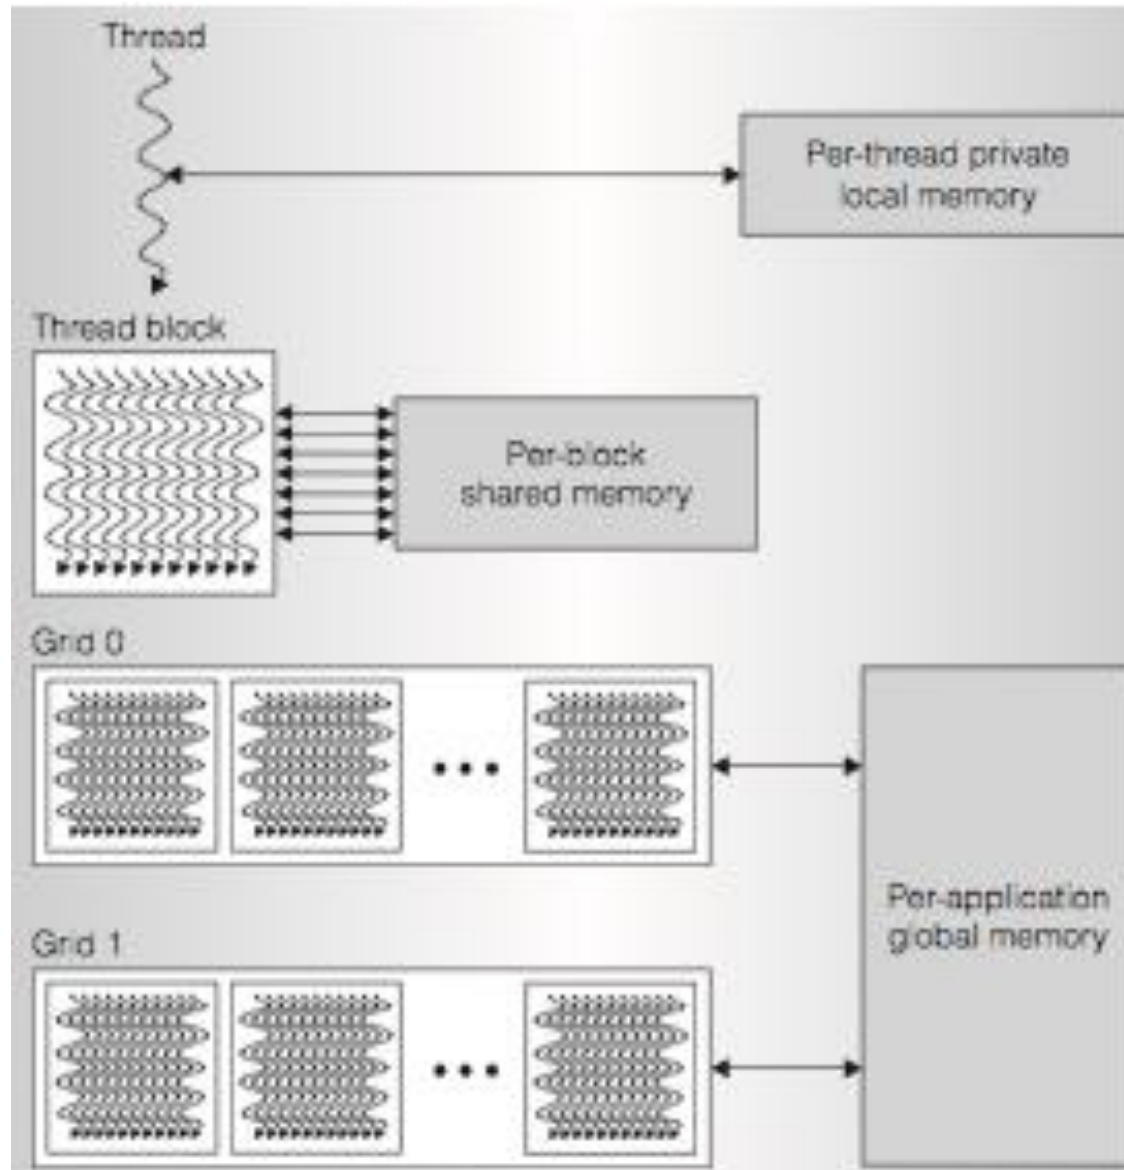


Tesla Products	Tesla K40	Tesla M40	Tesla P100
GPU	GK110 (Kepler)	GM200 (Maxwell)	GP100 (Pascal)
SMs	15	24	56
TPCs	15	24	28
FP32 CUDA Cores / SM	192	128	64
FP32 CUDA Cores / GPU	2880	3072	3584
FP64 CUDA Cores / SM	64	4	32
FP64 CUDA Cores / GPU	960	96	1792
Base Clock	745 MHz	948 MHz	1328 MHz
GPU Boost Clock	810/875 MHz	1114 MHz	1480 MHz
Peak FP32 GFLOPs <sup>1</sup>	5040	6840	10600
Peak FP64 GFLOPs <sup>1</sup>	1680	210	5300
Texture Units	240	192	224
Memory Interface	384-bit GDDR5	384-bit GDDR5	4096-bit HBM2
Memory Size	Up to 12 GB	Up to 24 GB	16 GB
L2 Cache Size	1536 KB	3072 KB	4096 KB
Register File Size / SM	256 KB	256 KB	256 KB
Register File Size / GPU	3840 KB	6144 KB	14336 KB
TDP	235 Watts	250 Watts	300 Watts
Transistors	7.1 billion	8 billion	15.3 billion
GPU Die Size	551 mm <sup>2</sup>	601 mm <sup>2</sup>	610 mm <sup>2</sup>
Manufacturing Process	28 nm	28 nm	16 nm FinFET

---

# **MEMORY MODEL OF CUDA AND COOPERATING THREADS**

# GPU Memory Hierarchy



[ Nvidia, 2010]



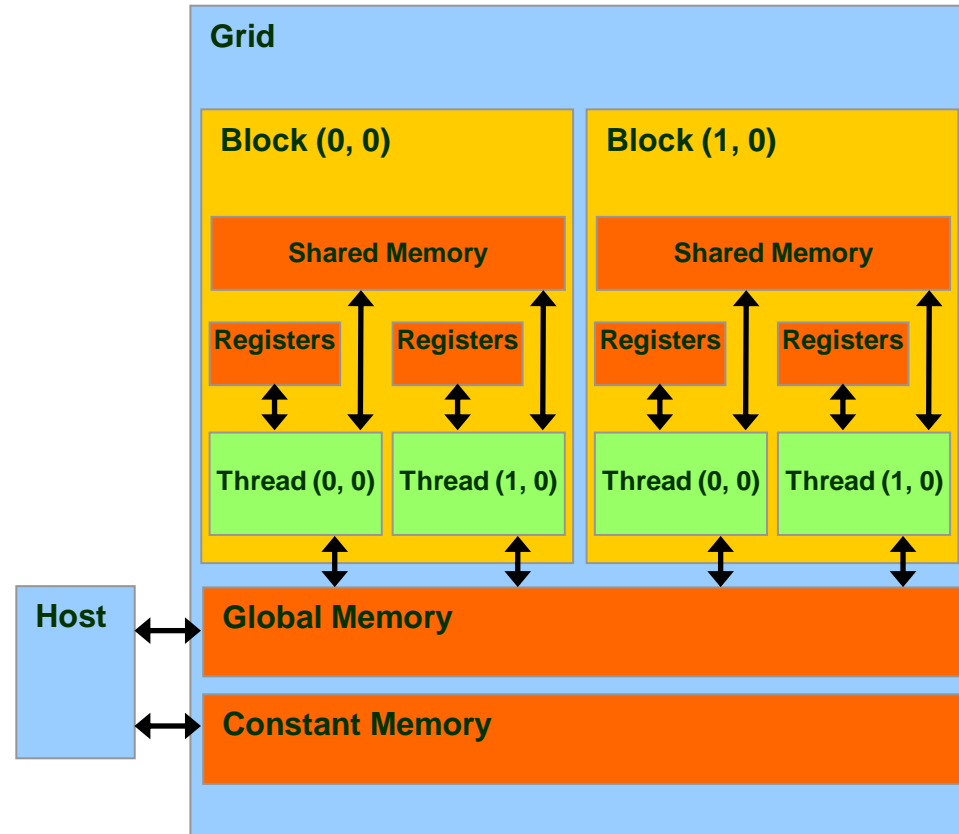
# CUDA Variable Type Qualifiers

Variable declaration	Memory	Scope	Lifetime
<code>__device__ __local__ int LocalVar;</code>	local	thread	thread
<code>__device__ __shared__ int SharedVar;</code>	shared	block	block
<code>__device__ int GlobalVar;</code>	global	grid	application
<code>__device__ __constant__ int ConstantVar;</code>	constant	grid	application

- ❑ `__device__` is optional when used with `__local__`, `__shared__`, or `__constant__`
- ❑ Automatic variables without any qualifier reside in a register
  - Except arrays that reside in local memory

# G80 Implementation of CUDA Memories


- ❑ Each thread can:
  - Read/write per-thread **registers**
  - Read/write per-thread **local memory**
  - Read/write per-block **shared memory**
  - Read/write per-grid **global memory**
  - Read/only per-grid **constant memory**



# 1D Stencil

---


- ❑ Consider applying a 1D stencil to a 1D array of elements
  - Each output element is the sum of input elements within a radius
- ❑ If radius is 3, then each output element is the sum of 7 input elements:





# Implementing Within a Block

- ❑ Each thread processes one output element
  - blockDim.x elements per block
- ❑ Input elements are read several times
  - With radius 3, each input element is read seven times




# Sharing Data Between Threads

---

- Terminology: within a block, threads share data via **shared memory**
- Extremely fast on-chip memory, user-managed
- Declare using `__shared__`, allocated per block
- Data is not visible to threads in other blocks

# Implementing With Shared Memory

- Cache data in shared memory
  - Read ( $\text{blockDim.x} + 2 * \text{radius}$ ) input elements from global memory to shared memory
  - Compute  $\text{blockDim.x}$  output elements
  - Write  $\text{blockDim.x}$  output elements to global memory
- Each block needs a **halo** of radius elements at each boundary



# Stencil Kernel

```
__global__ void stencil_1d(int *in, int *out) {  
    __shared__ int temp[BLOCK_SIZE + 2 * RADIUS];  
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;  
    int lindex = threadIdx.x + RADIUS;
```



```
    // Read input elements into shared memory
```

```
    temp[lindex] = in[gindex];  
    if (threadIdx.x < RADIUS) {  
        temp[lindex - RADIUS] = in[gindex - RADIUS];  
        temp[lindex + BLOCK_SIZE] =  
            in[gindex + BLOCK_SIZE];  
    }
```



# Stencil Kernel

---

```
// Apply the stencil
```

```
    result = 0;
```

```
    for (    offset = -RADIUS ; offset <= RADIUS ; offset++)
```

```
        result += temp[lindex + offset];
```

```
// Store the result
```

```
out[gindex] = result;
```


```
}
```

# Data Race!

---

- The stencil example will not work...
- Suppose thread 15 reads the halo before thread 0 has fetched it...


```
temp[lindex] = in[gindex];  
if (threadIdx.x < RADIUS) {  
    temp[lindex - RADIUS] = in[gindex - RADIUS];  
    temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];  
}
```

Store at temp[18] 

```
temp[lindex - RADIUS] = in[gindex - RADIUS];  
temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];  
}
```

Skipped, threadIdx > RADIUS

```
int result = 0;  
result += temp[lindex + 1];
```

Load from temp[19] 

## \_\_syncthreads()

- `__syncthreads () ;`
- Synchronizes all threads within a block
  - Used to prevent RAW / WAR / WAW hazards
- All threads must reach the barrier
  - In conditional code, the condition must be uniform across the block

# Stencil Kernel

---

```
__global__ void stencil_1d(int *in, int *out) {  
    int temp[BLOCK_SIZE + 2 * RADIUS];  
    int gindex = threadIdx.x + blockIdx.x * blockDim.x;  
    int lindex = threadIdx.x + radius;  
  
    // Read input elements into shared memory  
    temp[lindex] = in[gindex];  
    if (threadIdx.x < RADIUS) {  
        temp[lindex - RADIUS] = in[gindex - RADIUS];  
        temp[lindex + BLOCK_SIZE] = in[gindex + BLOCK_SIZE];  
    }  
  
    // Synchronize (ensure all the data is available)  
    syncthreads();  
}
```



# Stencil Kernel

---

```
// Apply the stencil
int result = 0;
for (int offset = -RADIUS ; offset <= RADIUS ; offset++)
    result += temp[lindex + offset];

// Store the result
out[gindex] = result;
}
```

# Review (1 of 2)

---

## ❑ Launching parallel threads

- Launch  $N$  blocks with  $M$  threads per block with `kernel<<<N,M>>> (...)` ;
- Use `blockIdx.x` to access block index within grid
- Use `threadIdx.x` to access thread index within block

## ❑ Allocate elements to threads:

```
int index = threadIdx.x + blockIdx.x * blockDim.x
```

## Review (2 of 2)

---

- ❑ Use `__shared__` to declare a variable/array in shared memory
  - Data is shared between threads in a block
  - Not visible to threads in other blocks
  
- ❑ Use `__syncthreads()` as a barrier
  - Use to prevent data hazards

---

# CUDA PERFORMANCE

---

Efficient data-  
parallel algorithms

+

Optimizations based  
on GPU Architecture

=

Maximum  
Performance

# Parallel Reduction


---

- Recall *Parallel Reduction* (sum)

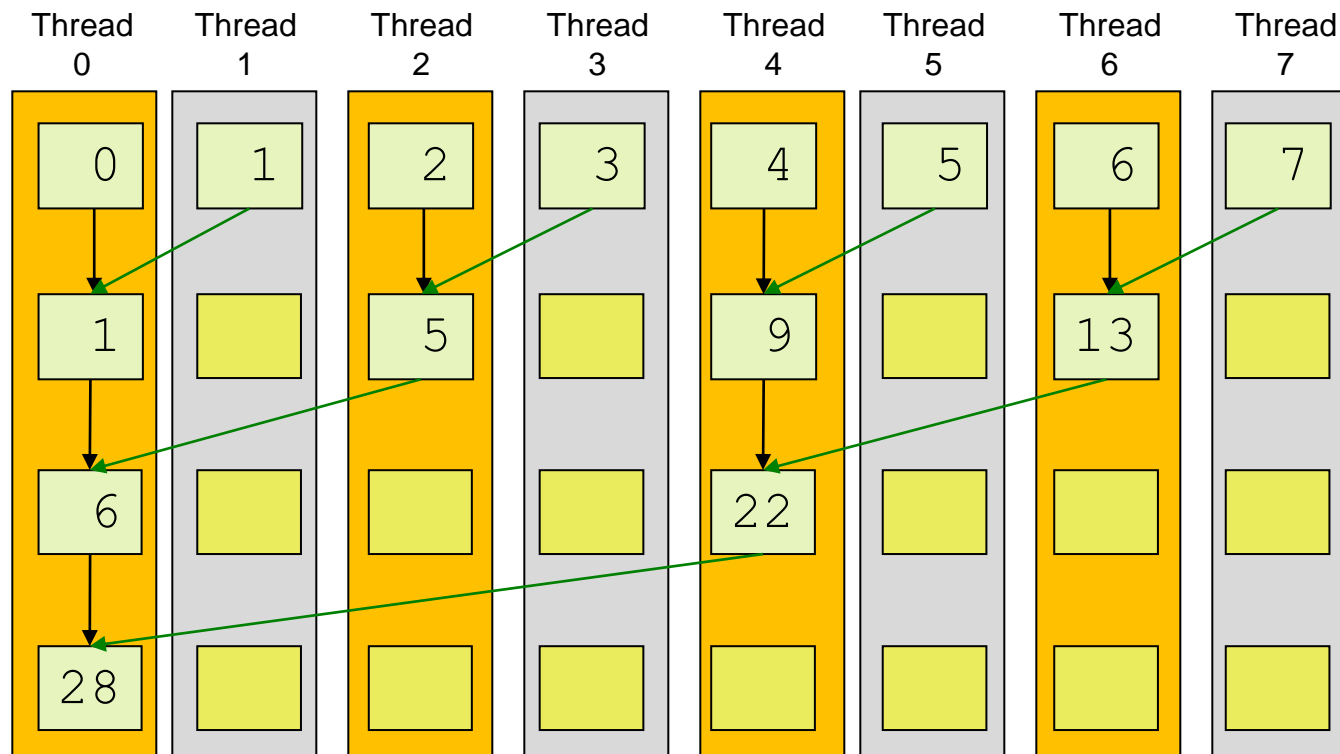


# Parallel Reduction

- ❑  $\log(n)$  passes for  $n$  elements
- ❑ How would you implement this in CUDA?



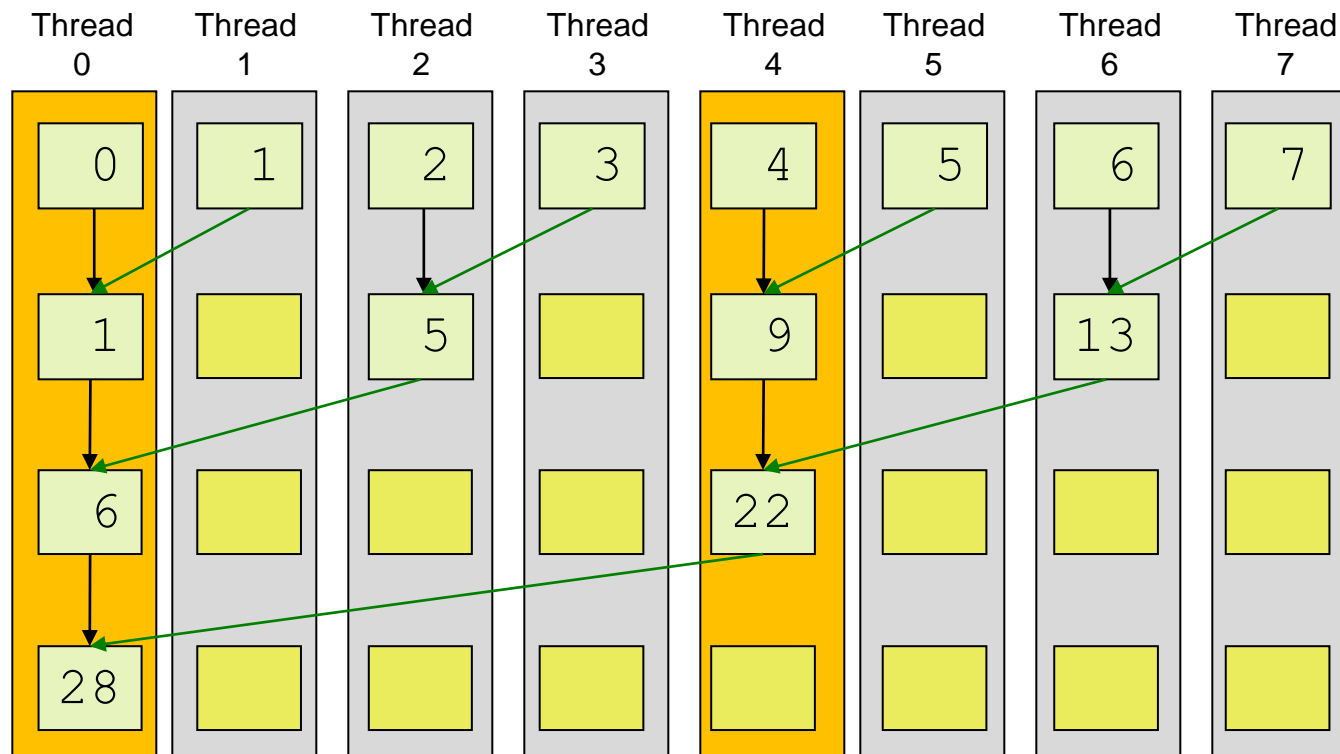
# Parallel Reduction



- ❑ 1<sup>st</sup> pass: threads 1, 3, 5, and 7 don't do anything
  - Really only need  $n/2$  threads for  $n$  elements

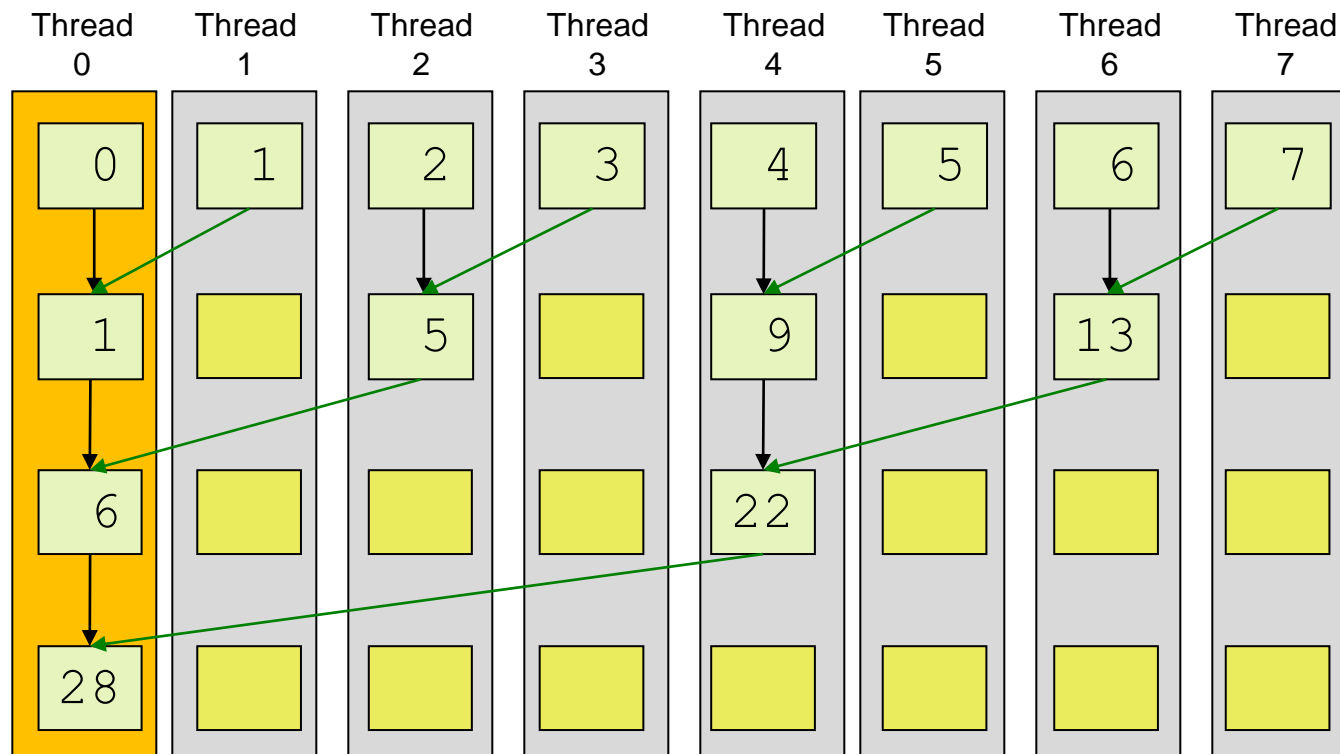


# Parallel Reduction




□ 2<sup>nd</sup> pass: threads 2 and 6 also don't do anything

# Parallel Reduction



□ 3<sup>rd</sup> pass: thread 4 also doesn't do anything

# Parallel Reduction



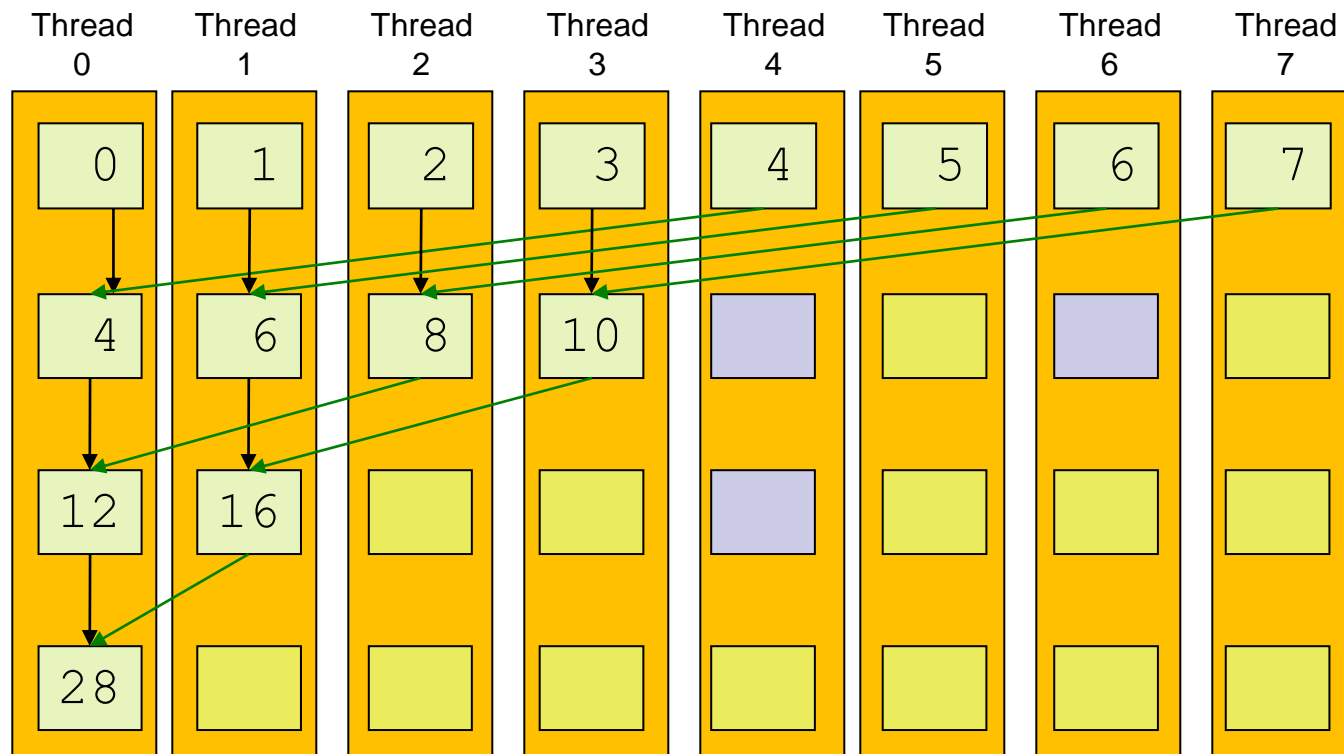
- ❑ In general, number of required threads cuts in half after each pass

# Parallel Reduction

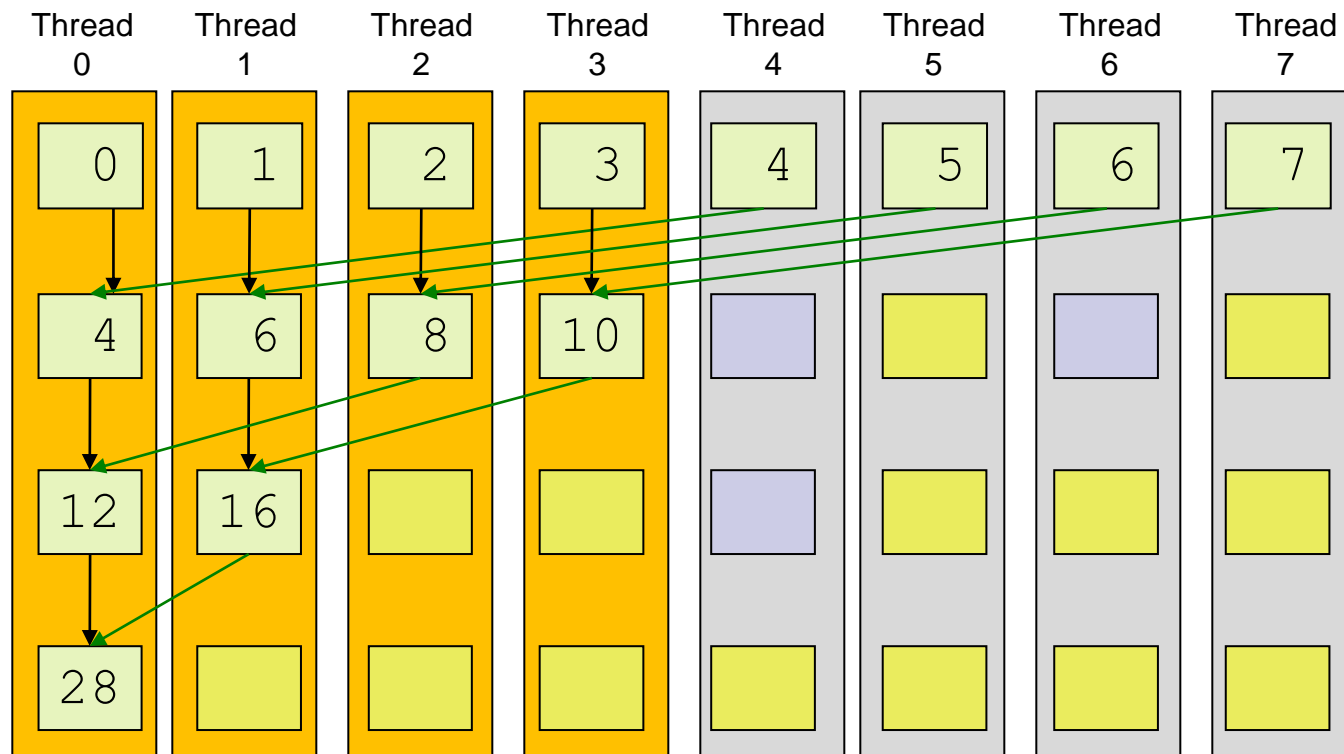
---

- What if we *tweaked* the implementation?

# Parallel Reduction

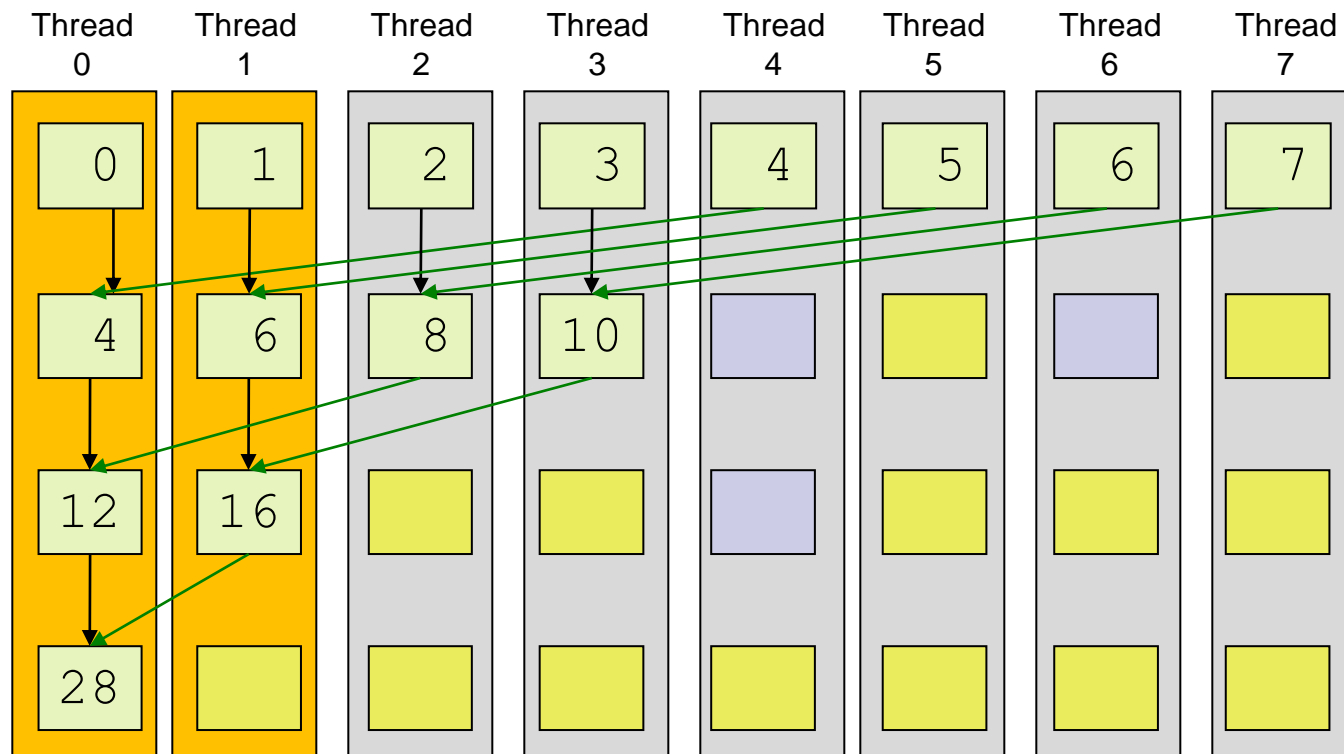


# Parallel Reduction



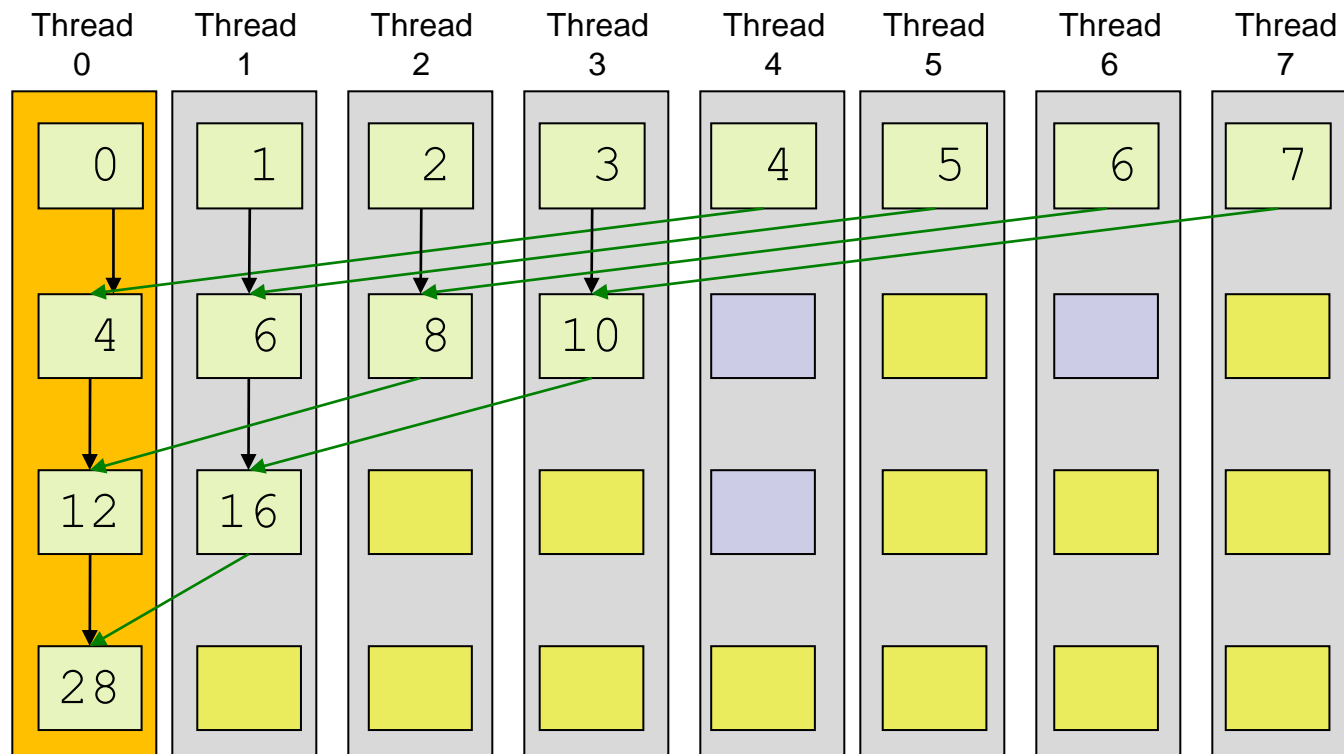
- ❑ 1<sup>st</sup> pass: threads 4, 5, 6, and 7 don't do anything
  - Really only need  $n/2$  threads for  $n$  elements

# Parallel Reduction



□ 2<sup>nd</sup> pass: threads 2 and 3 also don't do anything

# Parallel Reduction




□ 3<sup>rd</sup> pass: thread 1 also doesn't do anything




# Parallel Reduction

□ What is the difference?



stride = 1, 2, 4, ...



stride = 4, 2, 1, ...

# Parallel Reduction

---

□ What is the difference?

```
if (t % (2 * stride) == 0)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 1, 2, 4, ...

```
if (t < stride)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 4, 2, 1, ...

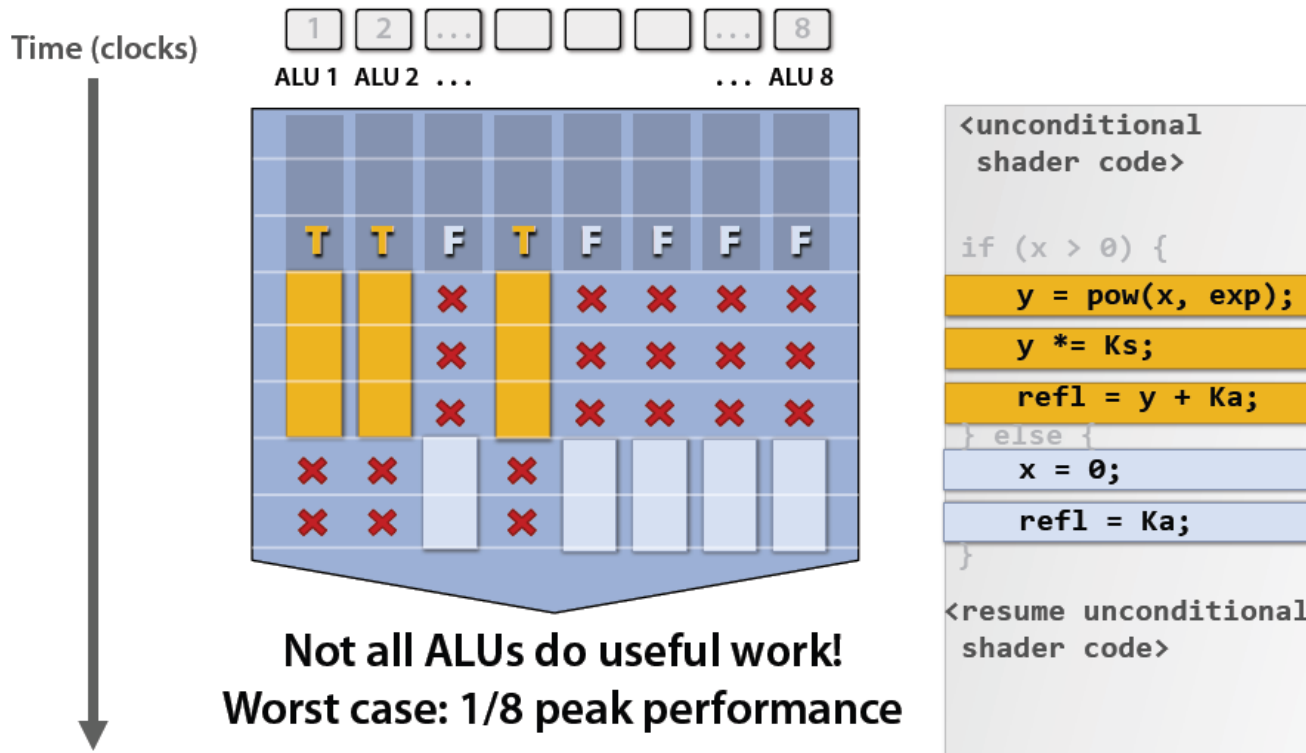
# Warp Partitioning

---

- ❑ *Warp Partitioning*: how threads from a block are divided into warps
- ❑ Knowledge of warp partitioning can be used to:
  - Minimize divergent branches
  - Retire warps early


# Warp Partitioning

*Divergent branches are within a warp!*




# Warp Partitioning

❑ Pretend `warpSize == 2`




`stride = 1, 2, 4, ...`




`stride = 4, 2, 1, ...`

# Warp Partitioning

## 1<sup>st</sup> Pass





stride = 1, 2, 4, ...



stride = 4, 2, 1, ...



# Warp Partitioning

## 2<sup>nd</sup> Pass



# Warp Partitioning


## 2<sup>nd</sup> Pass





# Warp Partitioning

## 2<sup>nd</sup> Pass



# Warp Partitioning

---

- ❑ Good partitioning also allows warps to be retired early.
  - Better hardware utilization

```
if (t % (2 * stride) == 0)
    partialSum[t] +=
        partialSum[t + stride];
```


stride = 1, 2, 4, ...

```
if (t < stride)
    partialSum[t] +=
        partialSum[t + stride];
```

stride = 4, 2, 1, ...


# Memory Coalescing (合并访存)

- Given a matrix stored *row-major* in *global memory*, what is a *thread*'s desirable access pattern?



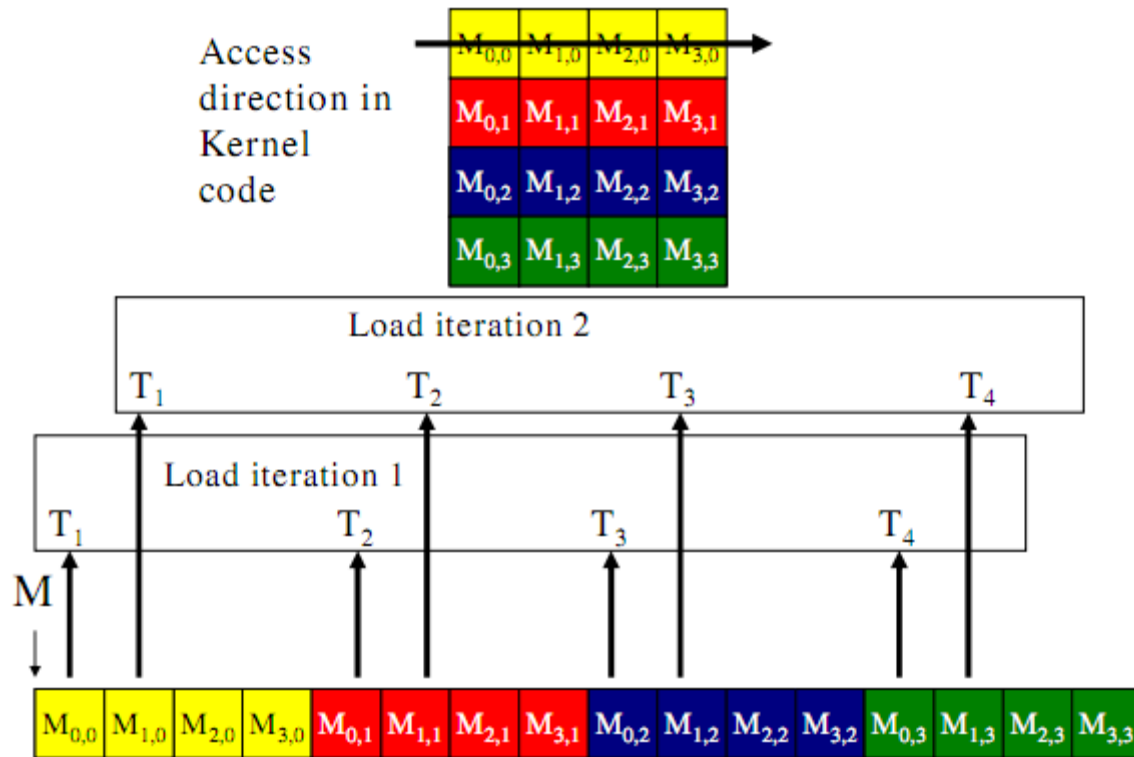
# Memory Coalescing

- Given a matrix stored *row-major* in *global memory*, what is a *thread*'s desirable access pattern?



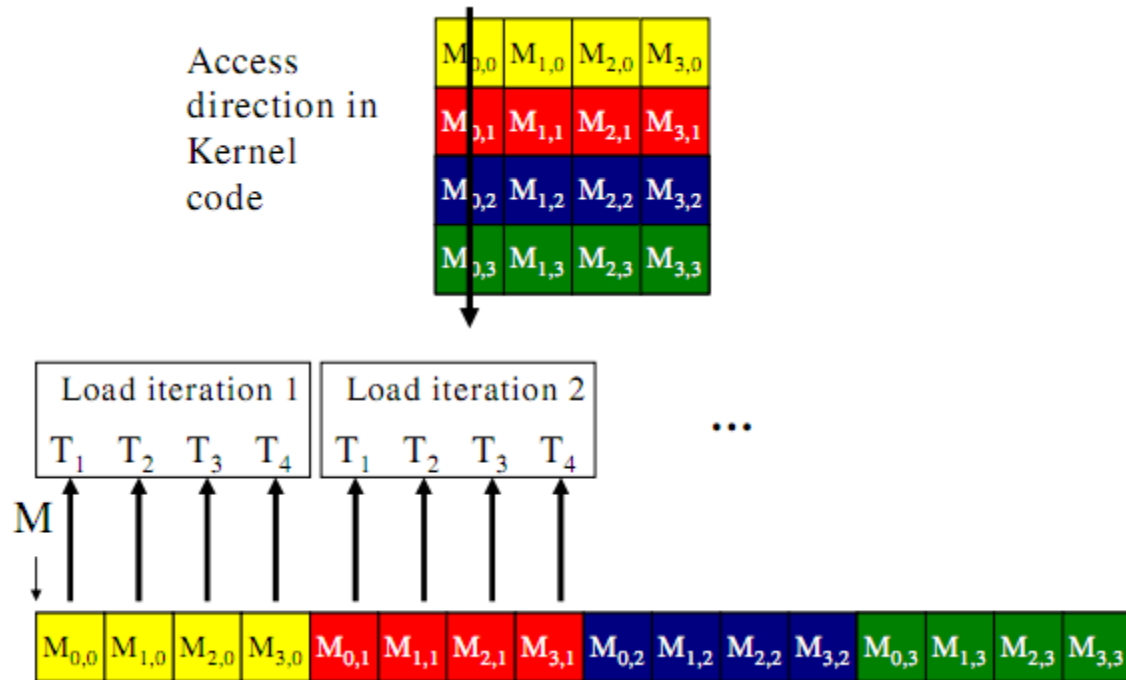
**a)** column after column? **b)** row after row?

# Memory Coalescing



a) column after column

# Memory Coalescing



b) row after row

# Memory Coalescing (合并访存)

---


- ❑ The GPU coalesce consecutive reads in a *half-warp* into a single read
- ❑ *Strategy*: read global memory in a coalesce-able fashion into shared memory
  - Then access shared memory randomly at maximum bandwidth
    - 📁 Ignoring *bank conflicts*...

# Bank Conflicts

---

## ❑ Shared Memory

- Sometimes called a *parallel data cache*
  - 📖 Multiple threads can access shared memory at the same time
- Memory is divided into *banks*






# Bank Conflicts


## ■ Bank Conflicts?

- Linear addressing  
stride == 1



## ■ Bank Conflicts?


- Random 1:1 Permutation



# Bank Conflicts


## ■ Bank Conflicts?

- Linear addressing  
stride == 2



## ■ Bank Conflicts?

- Linear addressing  
stride == 8




# Bank Conflicts

---

## □ Fast Path 1 (G80)


- All threads in a half-warp access different banks



# Bank Conflicts

## □ Fast Path 2 (G80)


- All threads in a half-warp access the same address



# Bank Conflicts

## ❑ Slow Path (G80)

- Multiple threads in a half-warp access the same bank
- Access is serialized
- What is the cost?



# Bank Conflicts

---


```
__shared__ float shared[256];  
// ...  
float f = shared[index + s * threadIdx.x];
```

□ For what values of  $s$  is this conflict free?

- Hint: The G80 has 16 banks

# Bank Conflicts

```
__shared__ float shared[256];  
// ...  
float f = shared[index + s * threadIdx.x];
```



no conflicts: stride 和 bank数目 没有公因子,  $s$ 必须为奇数

## 下一节

---

### □ 线程级并行性（TLP）

- 多核处理器中的关键问题：高速缓存一致性

From : H&P Computer Architecture: A Quantitative Approach,  
Fifth Edition, (5th edition)



# 高速缓存一致性与假共享： 例题

- ❑ 如下代码在SMP (shared memory multiprocessors) 环境下执行, sum和sum\_local是全局变量, 被NUM\_THREADS个线程所共享:

```
double sum=0.0, sum_local[NUM_THREADS];
```

```
#pragma omp parallel num_threads(NUM_THREADS)
```

```
//由NUM_THREADS个线程执行以下相同的代码段
```

```
{ int me = omp_get_thread_num();
```

```
  sum_local[me] = 0.0;
```

```
  #pragma omp for //并行for语句, 不同线程处理部分数据
```

```
  for (i = 0; i < N; i++)
```

```
    sum_local[me] += x[i] * y[i]; //将结果存入对应该线程的sum_local元素中
```

```
  #pragma omp atomic      //并行原子操作,
```

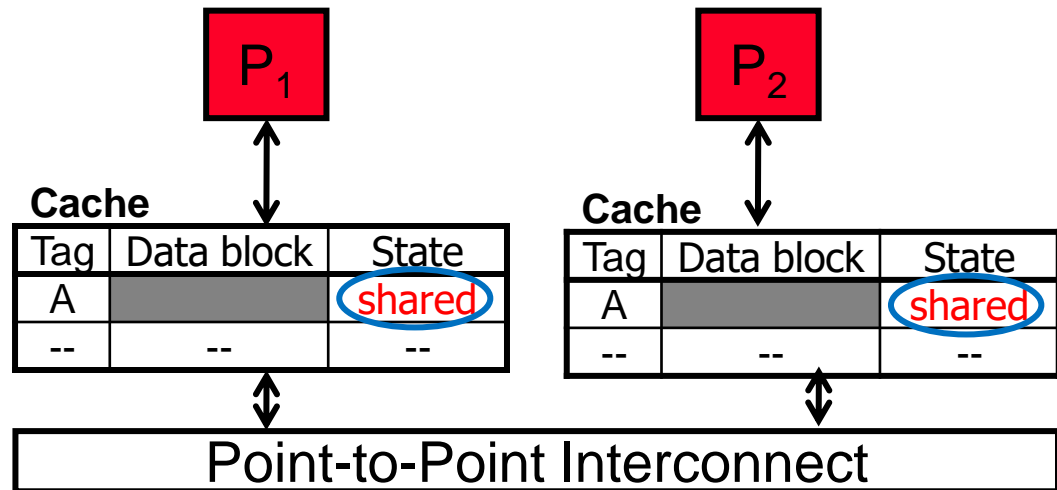
```
    sum += sum_local[me]; //求总和
```

```
}
```

# 高速缓存一致性和假共享

假设:

- **P1**写一个数据块内的第*i*个字
  - **P2**写同一块内的第*k*个字
- 会发生什么？



初始时，**P1**和**P2**共享一个数据块，  
私有**cache**中的状态都是**shared**

# 高速缓存一致性和假共享

- 高速缓存一致性协议以数据块为单位，而不是以字为单位
- 一个高速缓存数据块包含的字数多于1




一个高速缓存数据块


假共享:

- 当两个或更多处理器共享同一个数据块的不同部分时是假共享


# Bus Based Snooping Protocol



X is a shared variable that has a copy in all caches. Then a write occurred



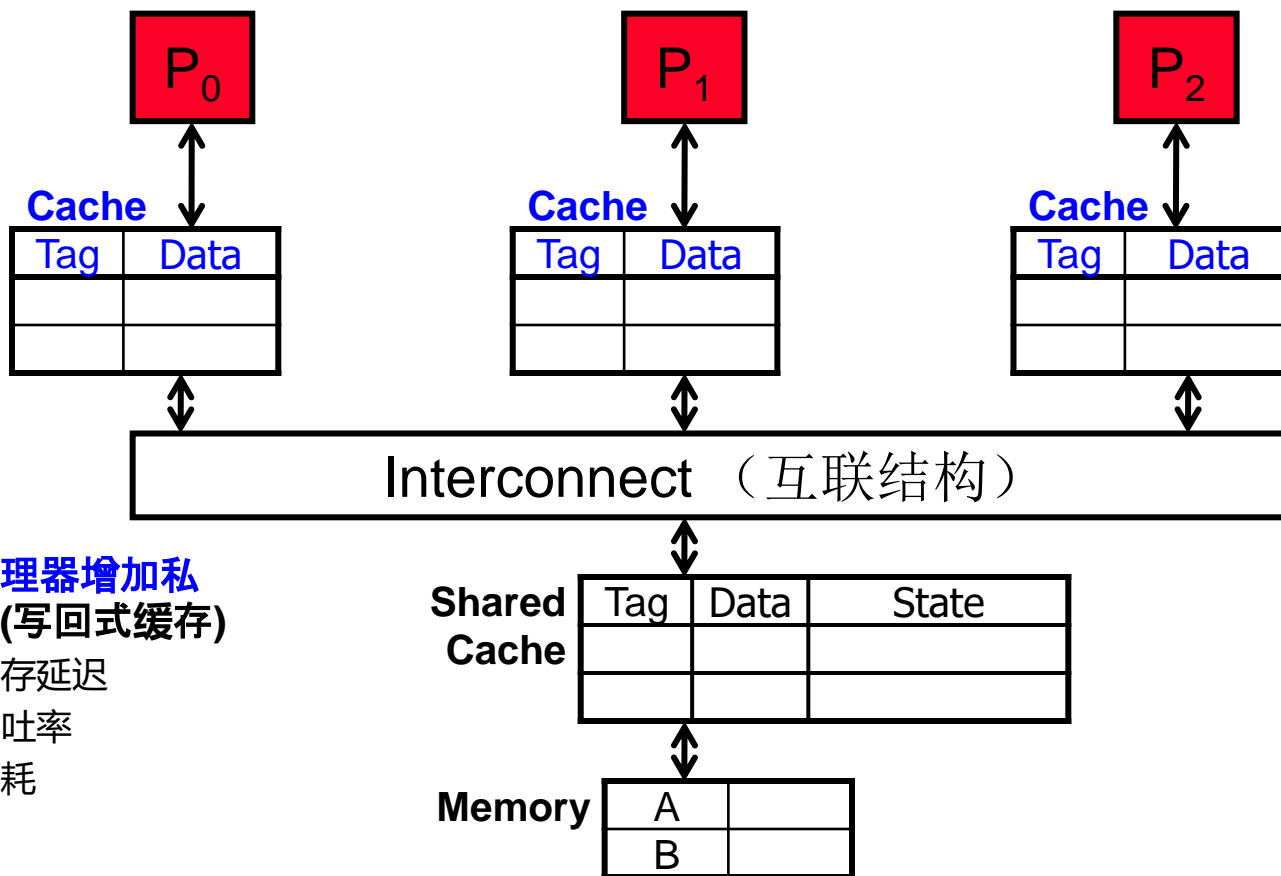
For **Write Invalidate** all the cache copies are marked as "invalid" except the most recent one



For **Write Update** all the cache copies are updated with the most recent value

Assume a write through cache protocol

# 增加私有高速缓存



为每一个处理器增加私有高速缓存(写回式缓存)

- 降低访存延迟
- 增加吞吐率
- 减少能耗



上海交通大学  
SHANGHAI JIAO TONG UNIVERSITY



谢 谢！

